



**University of Pisa**

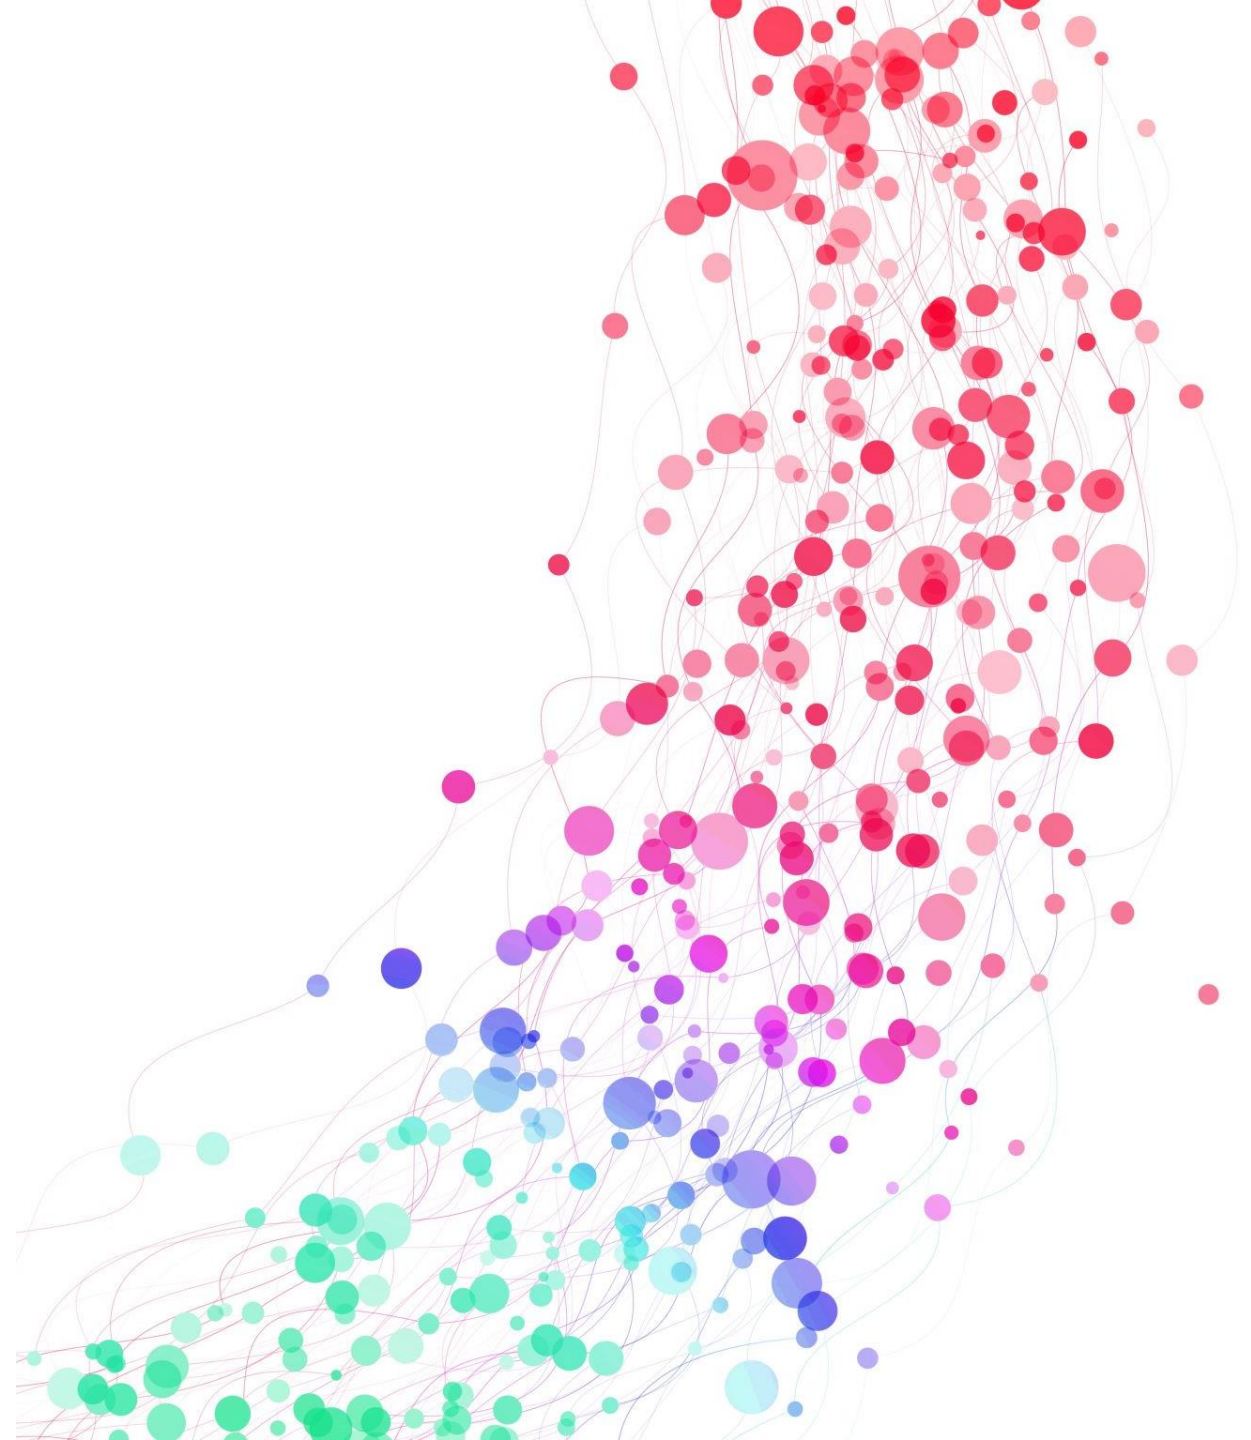
**Computer Architecture  
2023-24**

## **KMeans parallel implementation**

**Students**

**Carlo Pio Pace**

**Davide Vigna**



# The Algorithm

Input:

- Dataset of points
- Number of clusters  $K$
- Max # of iterations

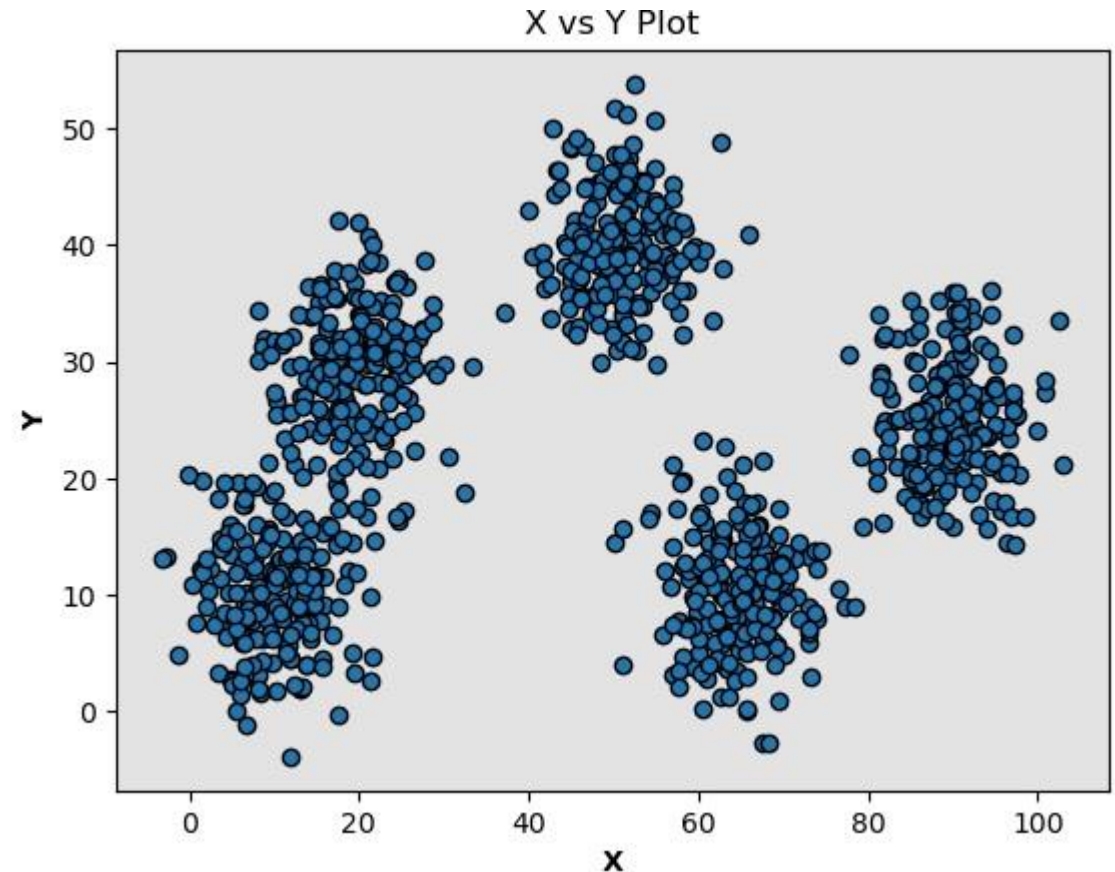
Output:

- A set of  $K$  points called Centroids
- The membership of each point in the dataset to one of the  $K$  clusters

Parameters:

- $K=5$
- Max # of iterations = 10

1) Start Execution



# The Algorithm

Input:

- Dataset of points
- Number of clusters  $K$
- Max # of iterations

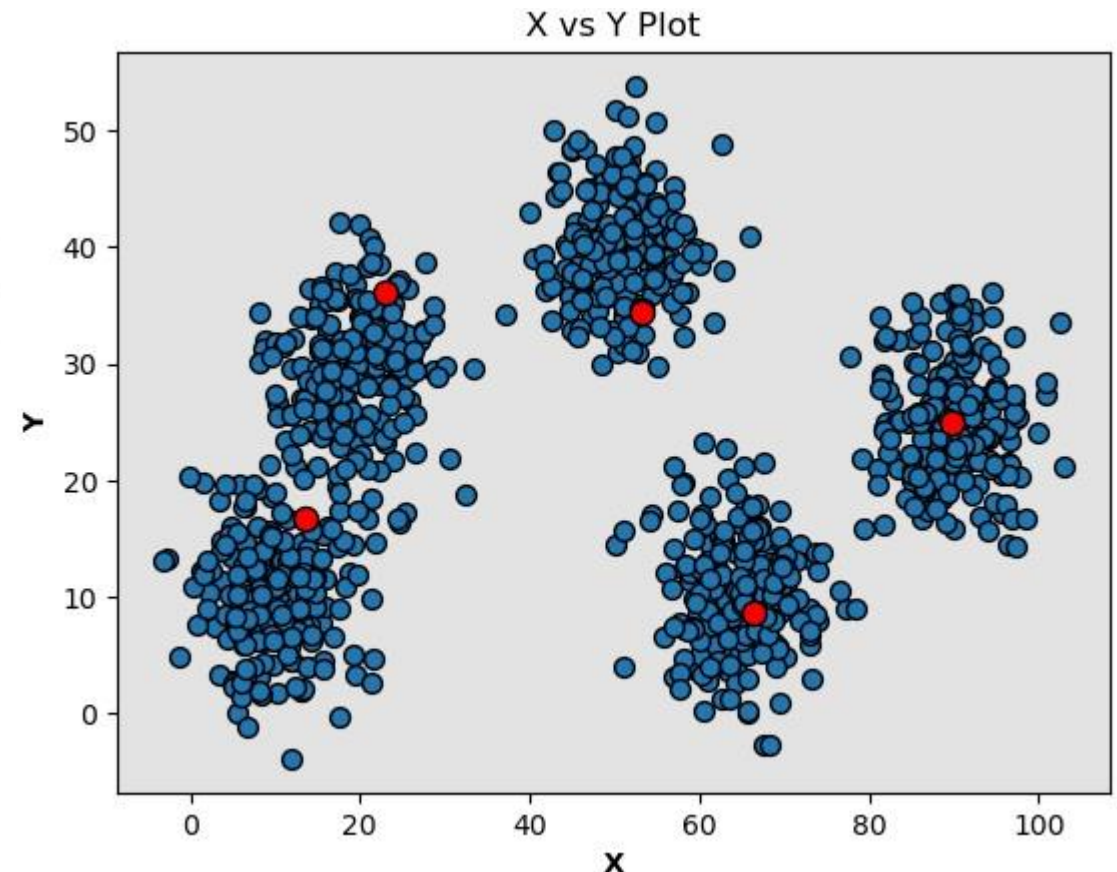
Output:

- A set of  $K$  points called Centroids
- The membership of each point in the dataset to one of the  $K$  clusters

Parameters:

- $K=5$
- Max # of iterations = 10

## 2) Initialization of $K$ random Centroids



# The Algorithm

Input:

- Dataset of points
- Number of clusters  $K$
- Max # of iterations

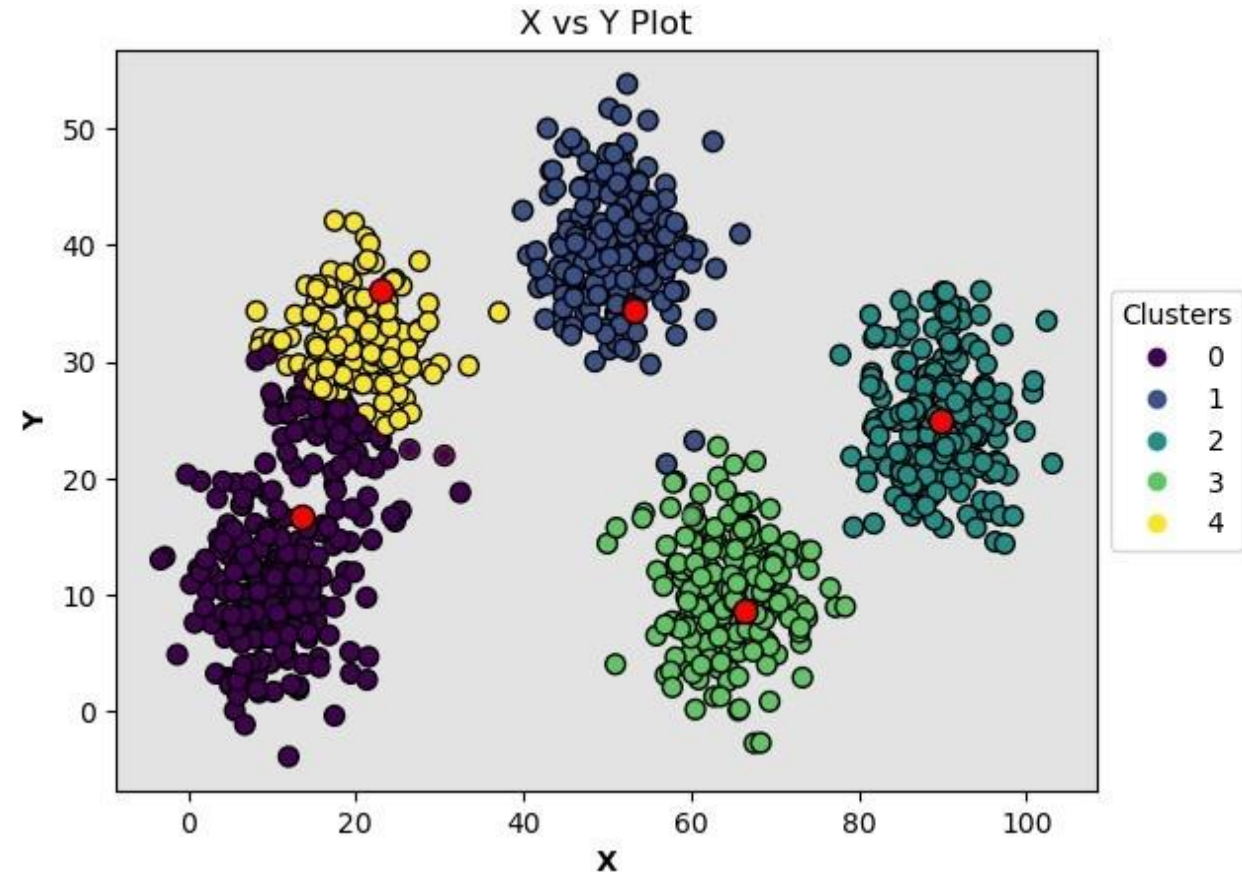
Output:

- A set of  $K$  points called Centroids
- The membership of each point in the dataset to one of the  $K$  clusters

Parameters:

- $K=5$
- Max # of iterations = 10

## 3) Assign points to clusters





# The Algorithm

Input:

- Dataset of points
- Number of clusters  $K$
- Max # of iterations

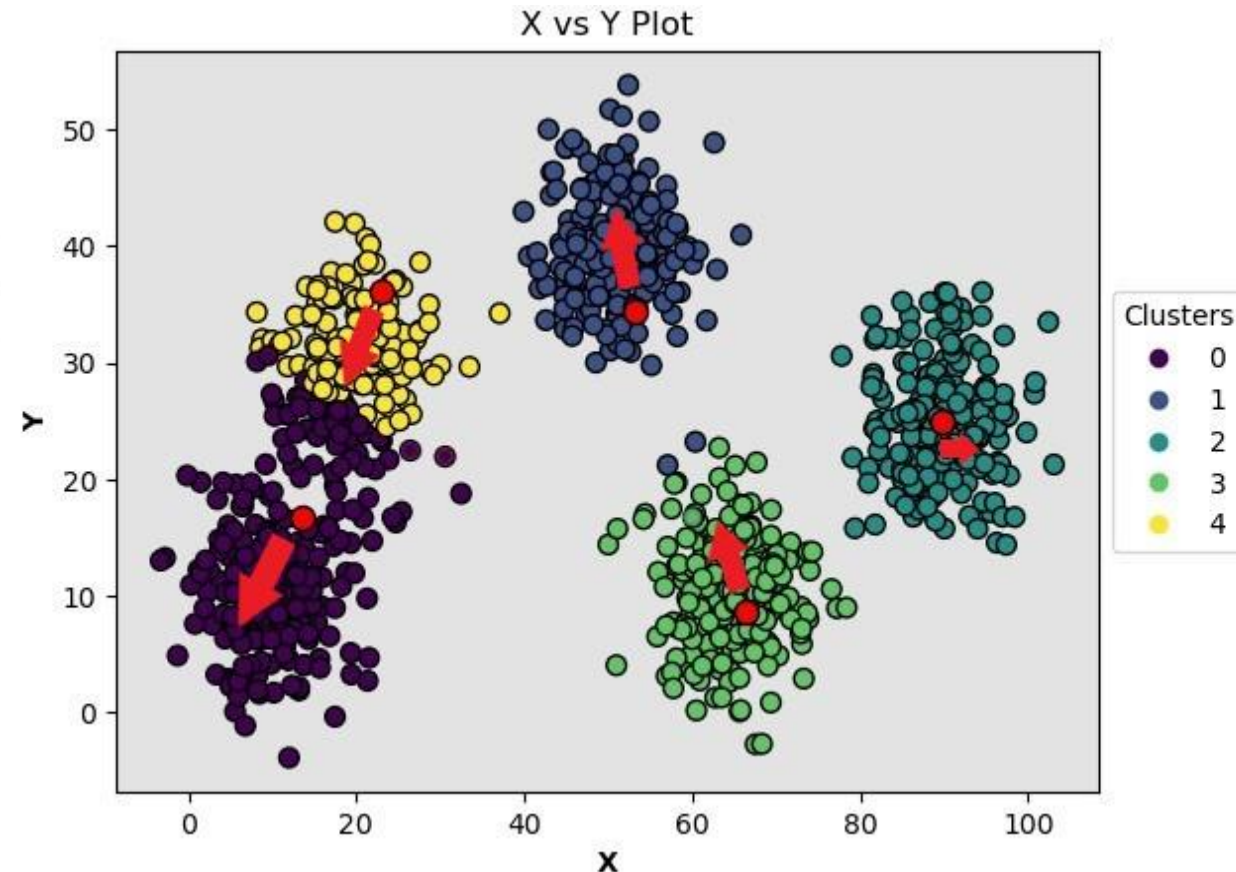
Output:

- A set of  $K$  points called Centroids
- The membership of each point in the dataset to one of the  $K$  clusters

Parameters:

- $K=5$
- Max # of iterations = 10

## 4) Update Centroids



# The Algorithm

Input:

- Dataset of points
- Number of clusters  $K$
- Max # of iterations

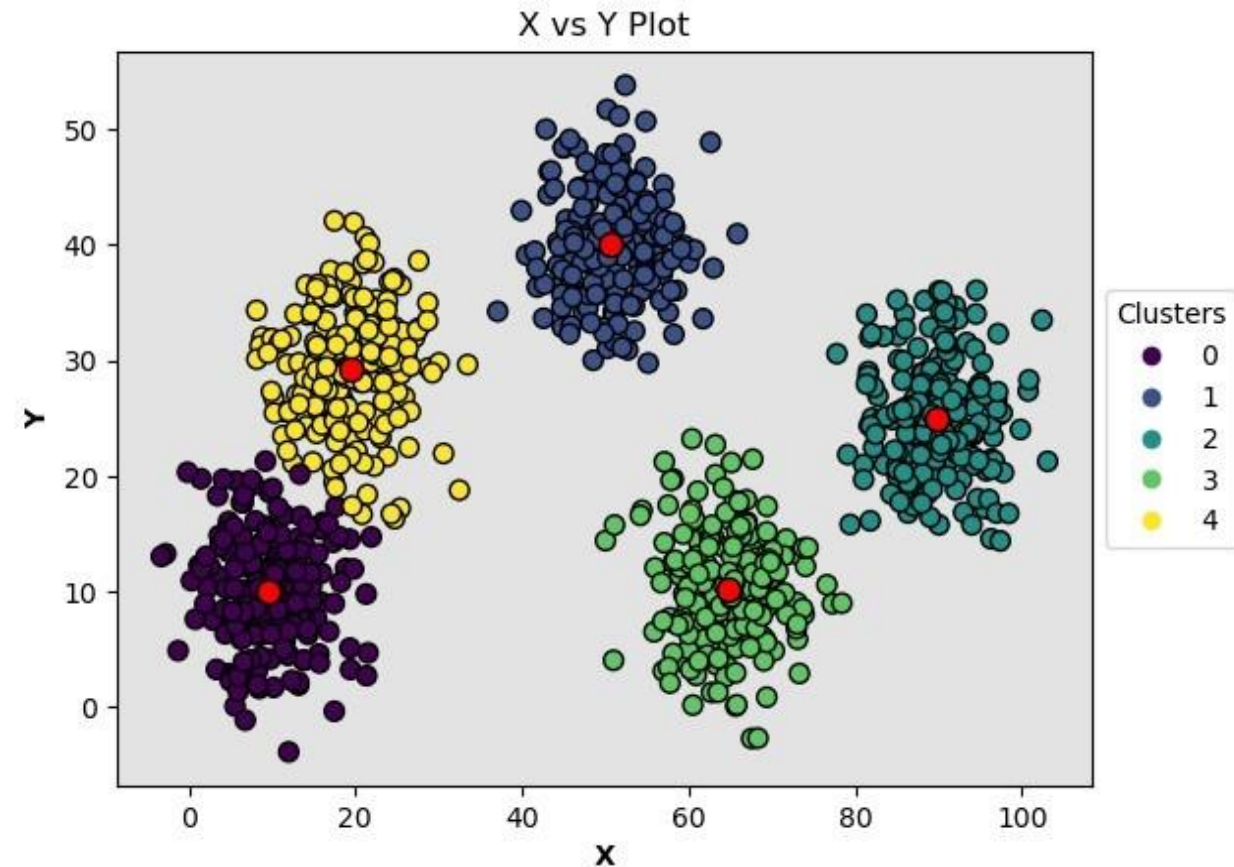
Output:

- A set of  $K$  points called Centroids
- The membership of each point in the dataset to one of the  $K$  clusters

Parameters:

- $K=5$
- Max # of iterations = 10

5) Repeat from 3 until Max # of iterations



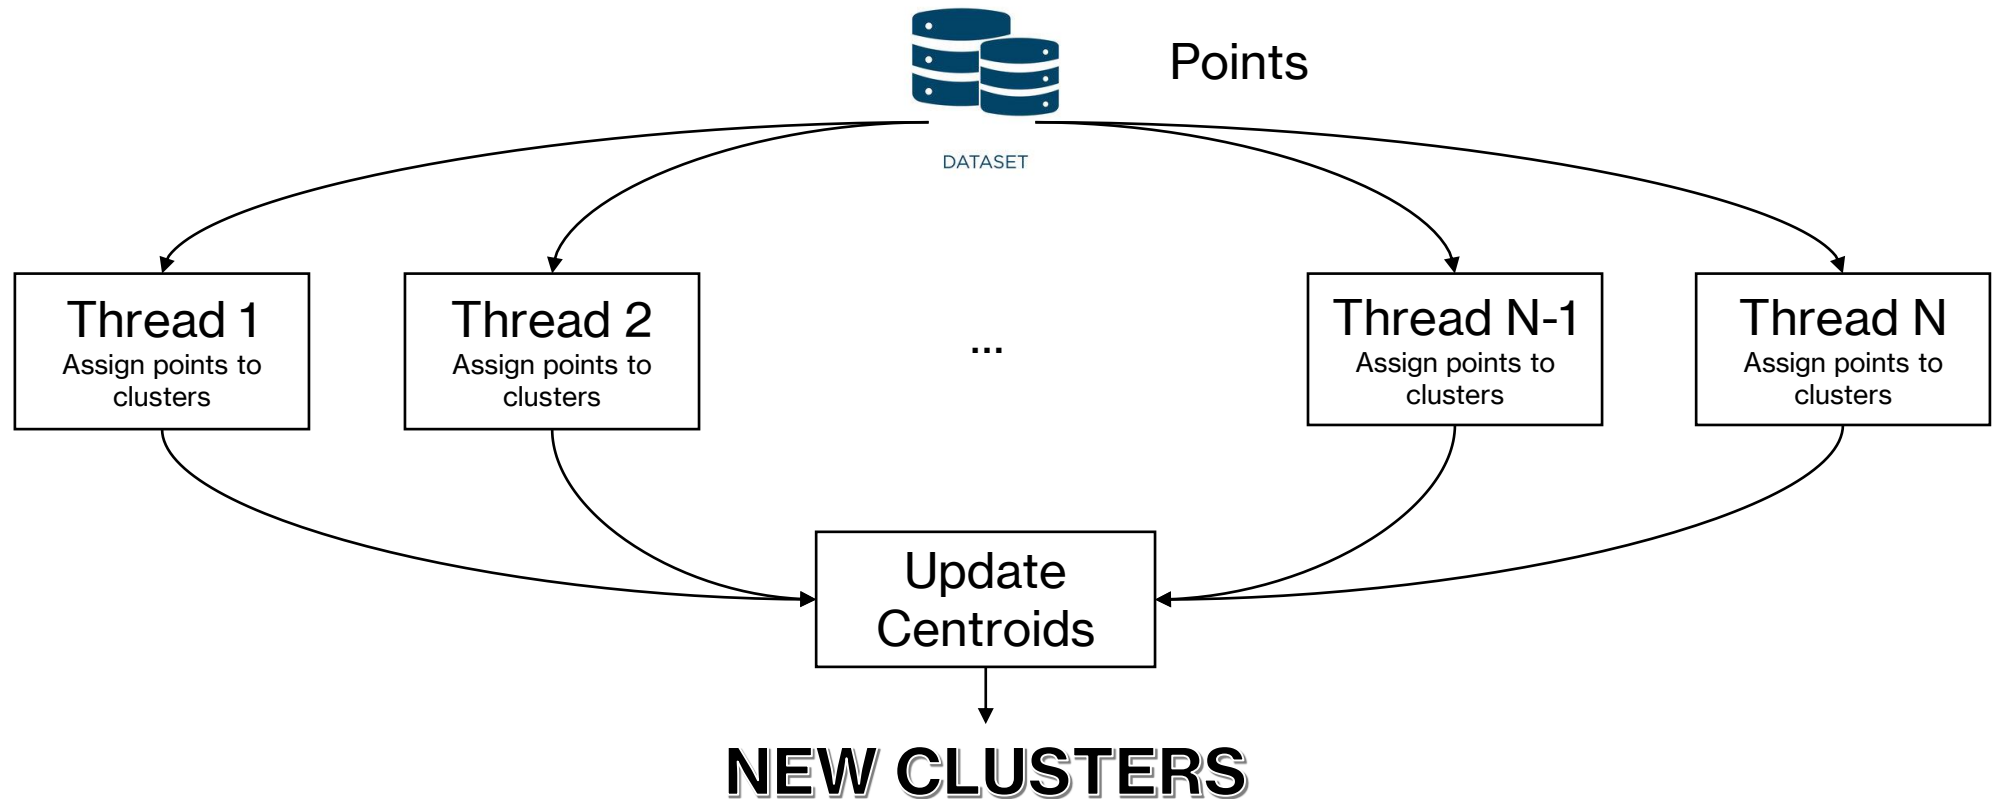


# The Goal

---

To obtain a solution that has an execution time of maximum 5 seconds on 10M dataset points.

# How to make it work in parallel





# How to make it work in parallel



30 (0,07%)	123
	124
28 (0,07%)	125
63 (0,16%)	126
	127
	128
120 (0,30%)	129
94 (0,23%)	130
28145 (70,22%)	131
	132
	133
3473 (8,66%)	134
	135

```
for (int indexOfPoint = arguments->startIndex; indexOfPoint < arguments->endIndex; indexOfPoint++) {
    Point *point=accessByIndex(arguments->points, indexOfPoint);

    for (int clusterIndex = 0; clusterIndex < arguments->K; clusterIndex++) {

        float sumPartial = 0;
        for (int dim = 0; dim < arguments->DIM; dim++) {
            Point* centroidPoint = accessByIndex(arguments->centroids, clusterIndex);
            sumPartial += pow(centroidPoint->coordinate[dim] - point->coordinate[dim], 2);
        }
        // Compute distance from a point to all centroids
        dists[clusterIndex] = sqrt(sumPartial);
    }
}
```

$$x \in R^{DIM}, y \in R^{DIM} \quad d(x, y) = \sqrt{\sum_{i=1}^{DIM} (x_i - y_i)^2}$$

Idea: exploit thread parallelism to work independently on different chunks of equal size of the dataset and then merge the results to update the centroids as in the sequential version.

# Architecture used

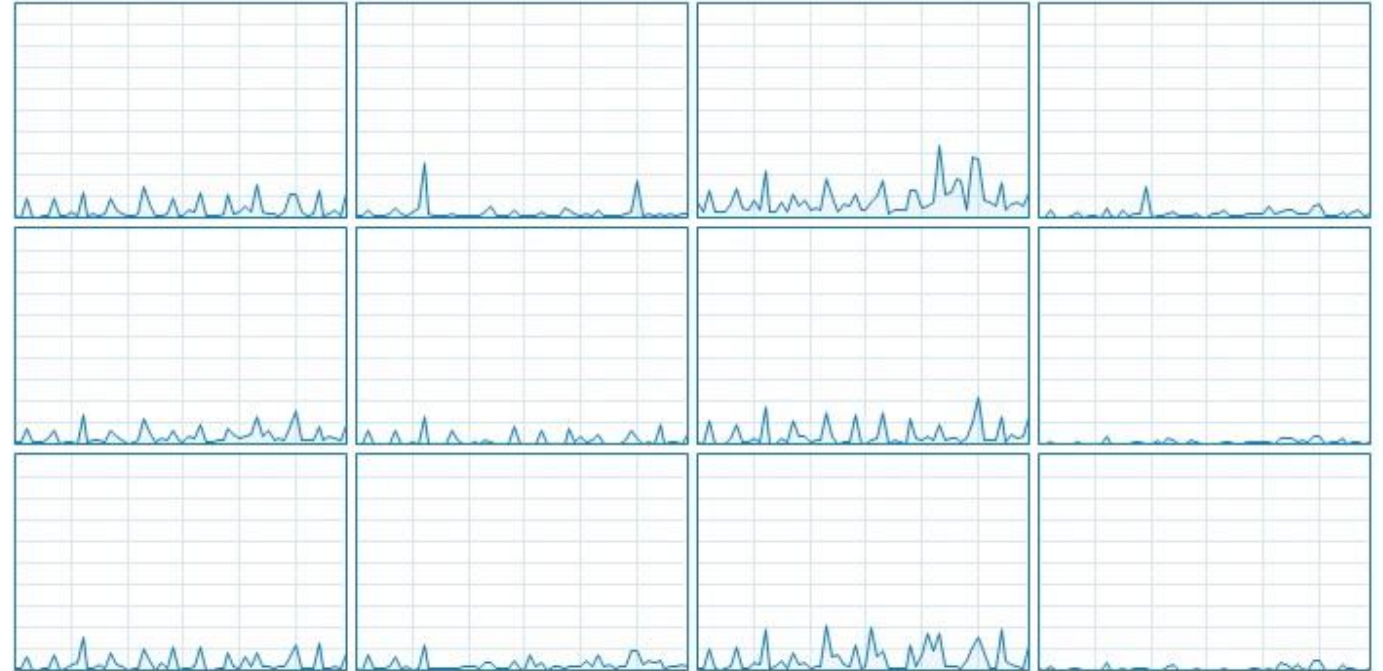
- CPU: Intel® Core™ i7-8700 Processor
  - 6 Cores, 12 Threads
  - Clock: Up to 3.20 Ghz
  - Max Turbo Frequency: 4.60 GHz
  - Cache:
    - Level 1: 384 KB
    - Level 2: 1.5 MB
    - Level 3: 12 MB
- RAM: 16 GB DDR4 2133 MHz

## CPU

Intel(R) Core(TM) i7-8700 CPU @ 3.20GHz

% di utilizzo in 60 secondi

100%



Utilizzo	Velocità		Velocità di base:	3,19 GHz
7%	3,16 GHz		Processori fisici:	1
			Cores:	6
Processi	Thread	Handle	Processori logici:	12
263	3828	143324	Virtualizzazione:	Abilitato
Tempo di attività			Cache L1:	384 KB
1:12:10:18			Cache L2:	1,5 MB
			Cache L3:	12,0 MB

## Architecture used

- CPU: Intel® Core™ i7-8700 Processor
  - 6 Cores, 12 Threads
  - Clock: Up to 3.20 Ghz
  - Max Turbo Frequency: 4.60 GHz
  - Cache:
    - Level 1: 384 KB
    - Level 2: 1.5 MB
    - Level 3: 12 MB
- RAM: 16 GB DDR4 2133 MHz

## Cache details

- Level 1 :
  - Data : 32 Kbytes 8-way set associative;
  - Instructions : 32 Kbytes 8-way set associative;
- Level 2 : 256 Kbytes 4-way set associative;
- Level 3: 12 Mbytes 16-way set associative;

# Measuring method

Function `clock()` in `time.h` C library

Why:

- It calculates CPU time and not Wall-clock time (with small errors)
- On modern machines, precision of ms
- C is a standard
- C is the language in which we developed the algorithm

How it works:

The function `clock()` defined in this library returns the corresponding cpu time at the instant when it's called.

# Measuring method

```
@echo off
REM parameter 1 Number of algorithm execution
REM parameter 2 Dataset size
REM parameter 3 Maximum number of thread to reach

REM Check if all three parameters are provided
if "%~3"==" " (
    echo Please provide three input parameters.
    exit /b 1
)

set "exe_path=filepath"
set "result_file=results.txt"

for /l %%i in (1, 1, %3) do (
    (
        echo datasetPath=clustering_dataset_%2.csv
        echo csvCharSplit=,
        echo numberOfClusters_K=5
        echo maxIterations=10
        echo nrThreads=%%i
    ) > config.properties

    for /L %%c in (1,1,%%1) do (
        echo Running .exe program for thread %%i
        %exe_path%
        echo Program executed for thread %%i.
    )

    REM Rename the file to the desired format
    rename %result_file% KmeansMeasure_%2_%%i
)
pause
```

```
start_alg = clock();

for (int nrIteration = 0; nrIteration < MAX_ITERATIONS; nrIteration++) {

    // Wake up all the suspended threads.
    for (int i = 0; i < NR_THREADS; i++) {
        ResumeThread(handles[i]);
    }

    // Wait until all the threads has produced their results and then come to sleep.
    WaitForMultipleObjects(NR_THREADS, suspendEvents, TRUE, INFINITE);

    // Aggregate partial results
    for (int j = 0; j < K; j++) {
        Point* totSum = accessByIndex(&sums, j);
        for (int i = 0; i < NR_THREADS; i++) {

            Point* partialSum = accessByIndex(results[i].sums, j);

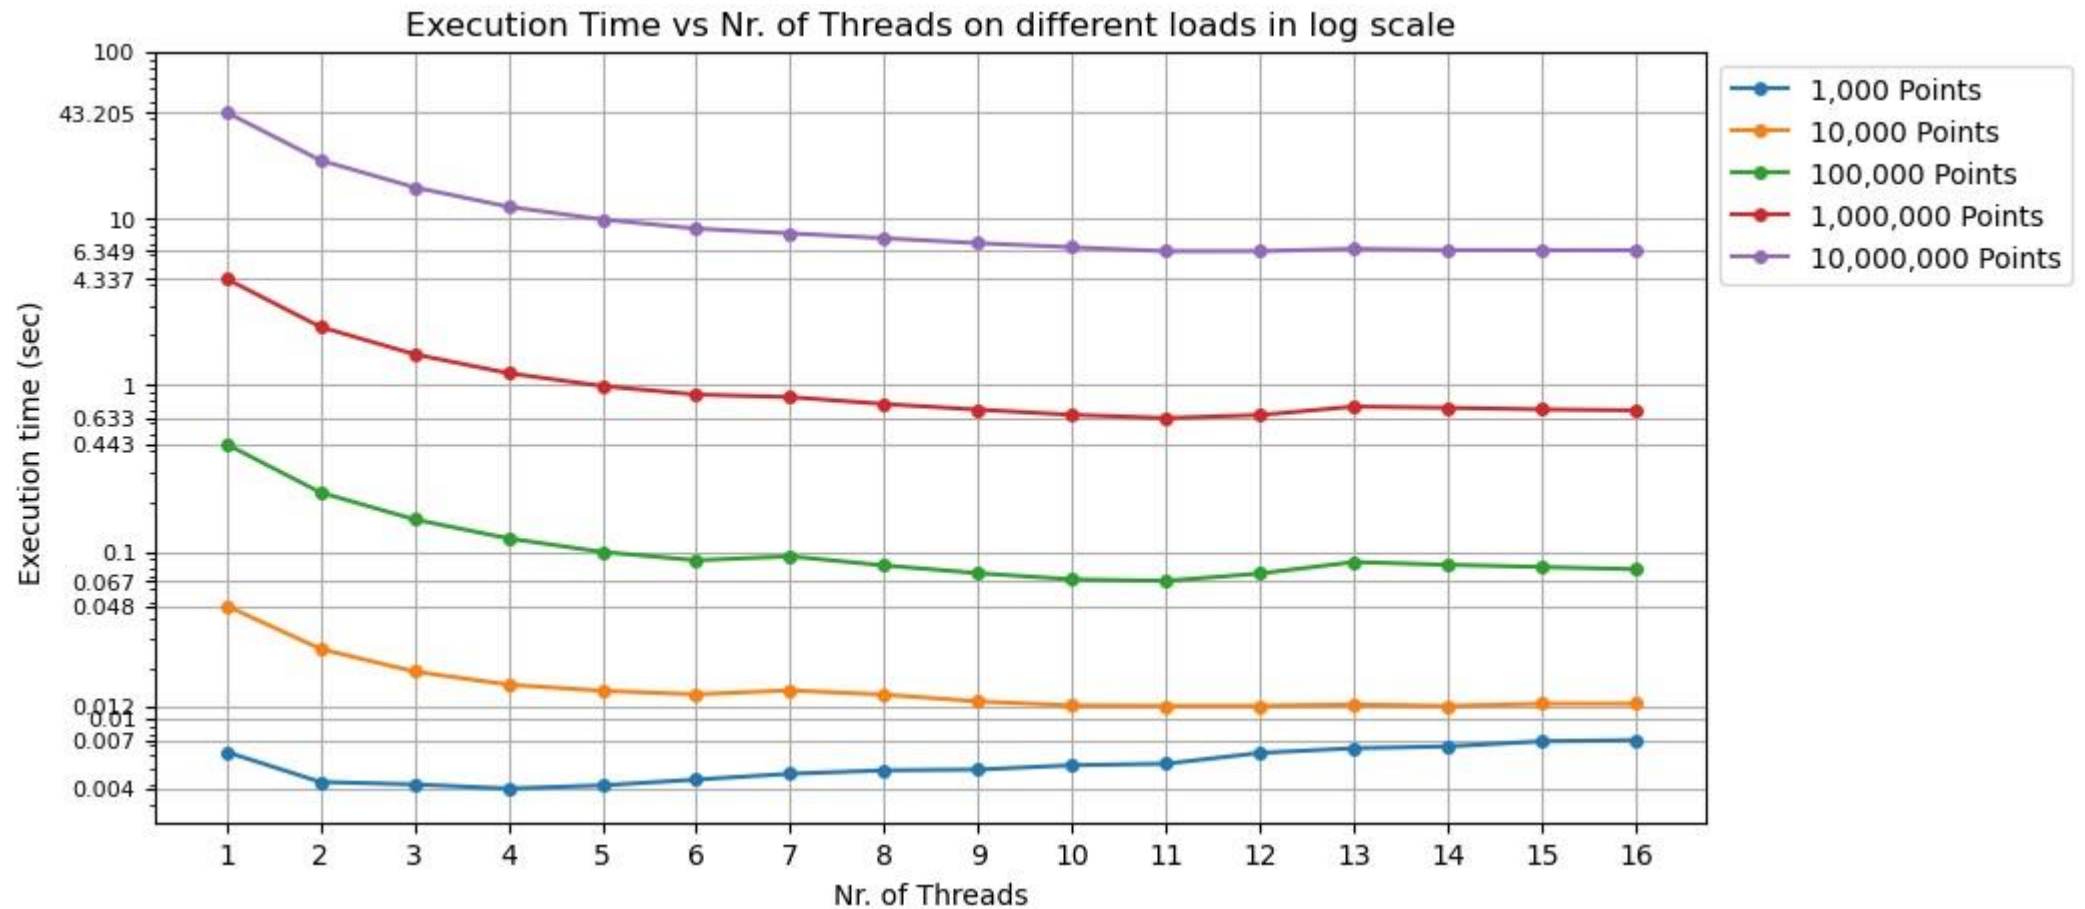
            for (int dim = 0; dim < DIM; dim++) {
                totSum->coordinate[dim] += partialSum->coordinate[dim];
            }
            counts[j] += results[i].counts[j];
        }
    }
    updateCentroids(K, DIM, &centroids, &sums, counts);

    // Reset the variables to track suspension next iterations.
    for (int i = 0; i < NR_THREADS; i++) {
        ResetEvent(suspendEvents[i]);
    }
}

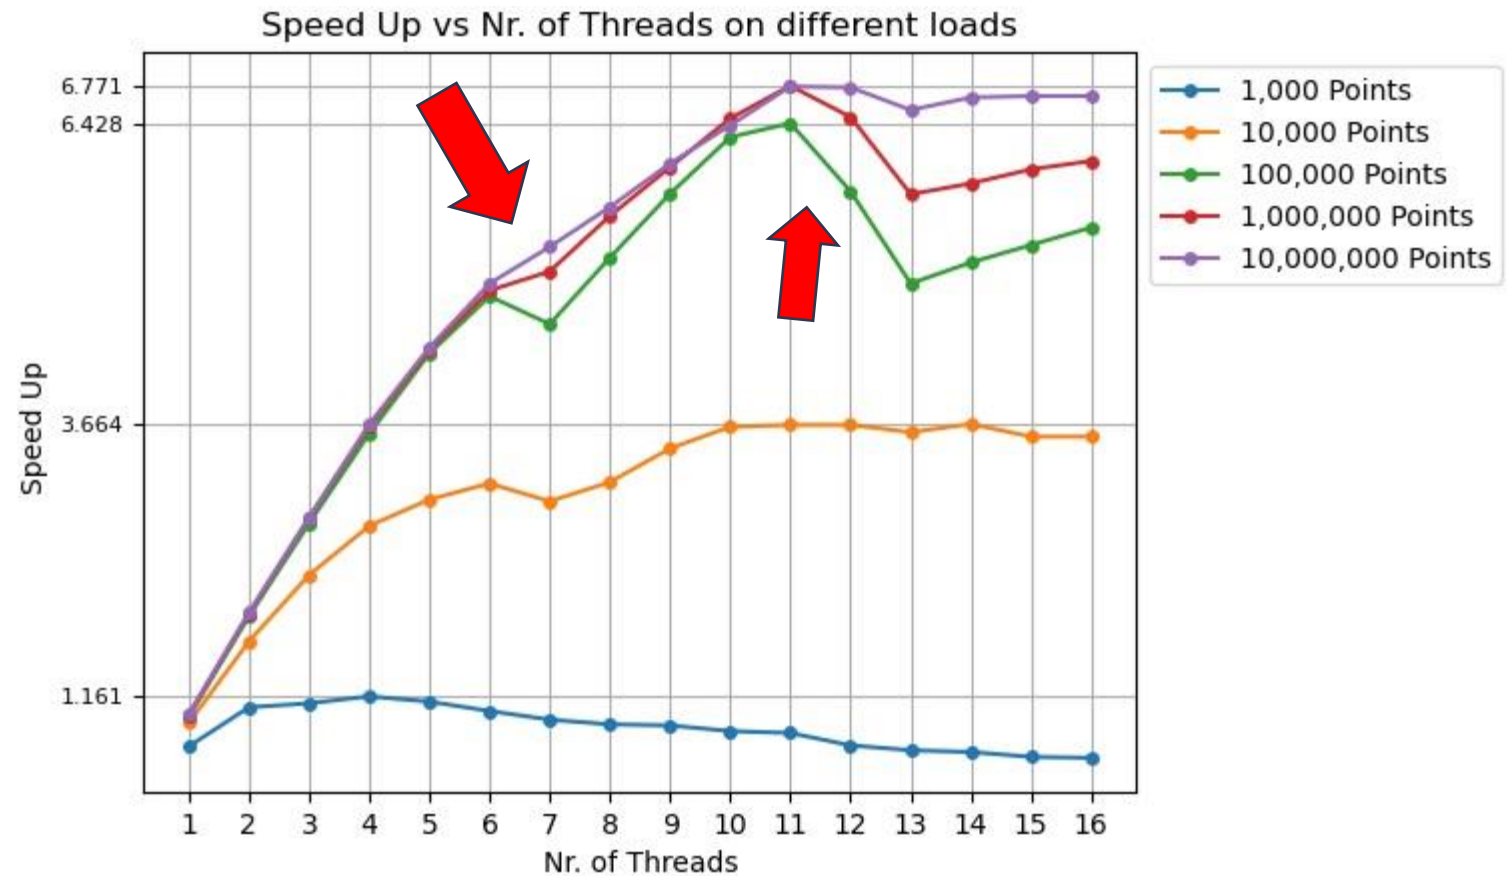
end_alg = clock();
```



# CPU time comparison



# Speed Up



# Speed Up – What's Up?



CORE 0	CPU 0	CPU 1
CORE 1	CPU 2	CPU 3
CORE 2	CPU 4	CPU 5
CORE 3	CPU 6	CPU 7
CORE 4	CPU 8	CPU 9
CORE 5	CPU 10	CPU 11

How the software threads are assigned to our logical core units?

The O.S. schedule them according his own policy

# Speed Up – What's Up?

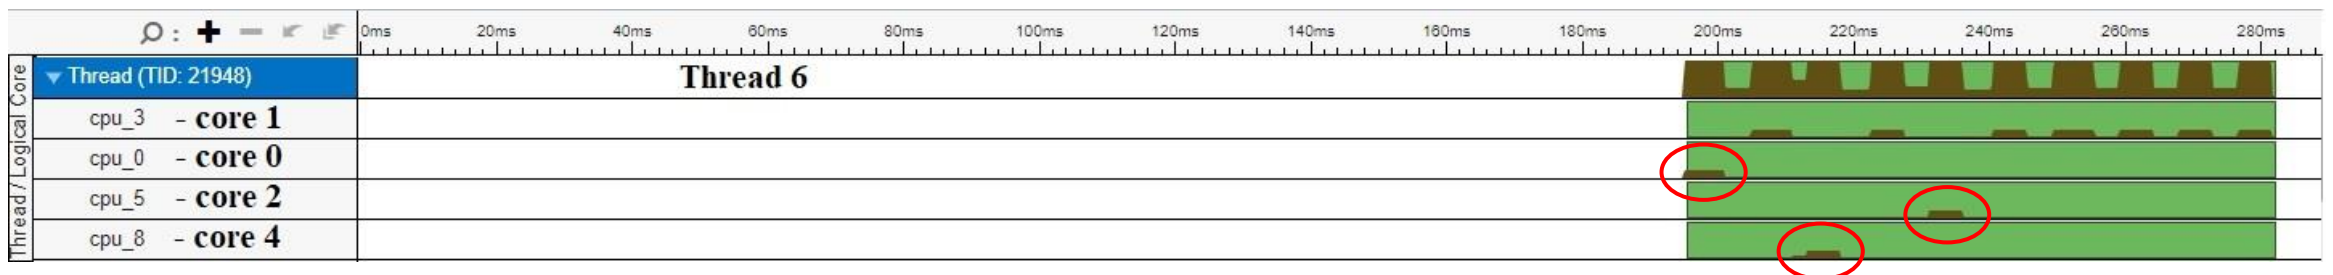
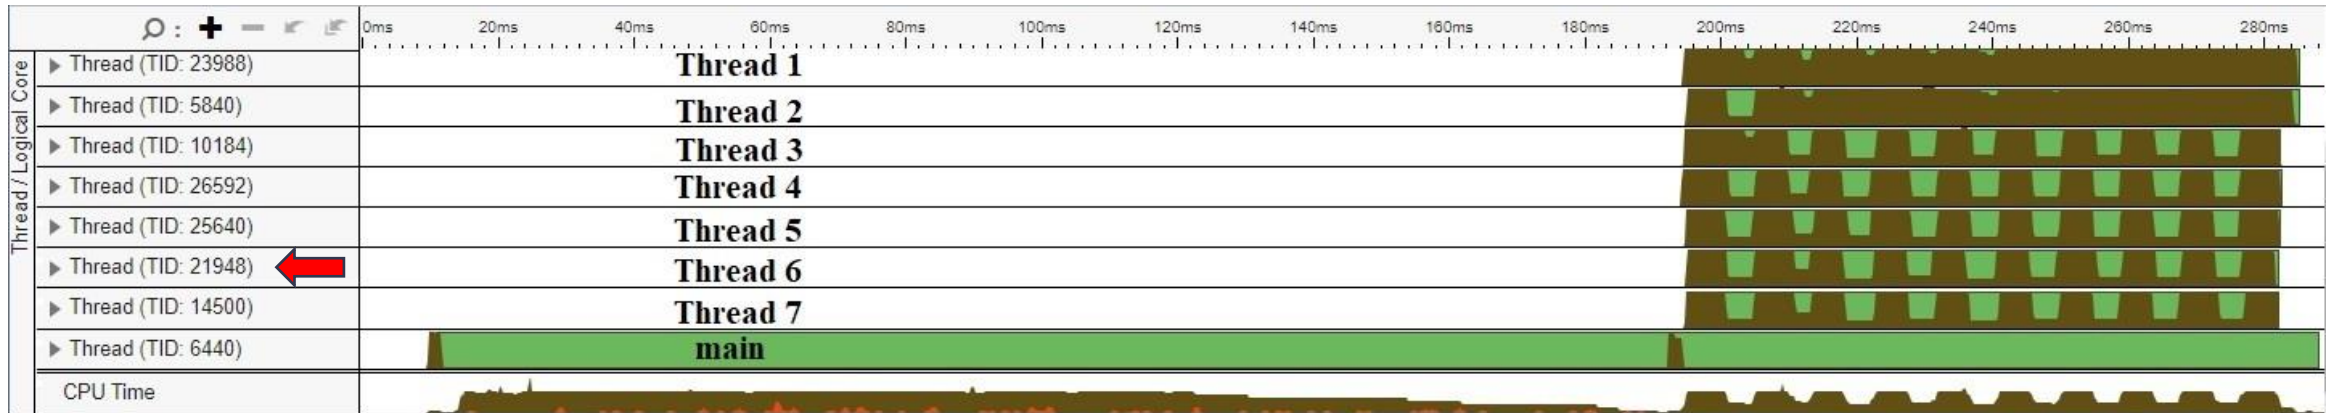


- 100 000 Points dataset

- 7 Threads

- Clusters=5 & Max # of iterations = 10

● Running ● CPU in use



# Speed Up – What's Up?

## Affinity:





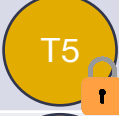
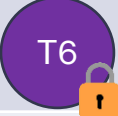
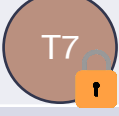
The explicit request to O.S. to assign a process or thread to a specific logical core inside a specific physical core.



# Using affinity



# Following logical cores order

CORE 0	CPU 0		CPU 1	
CORE 1	CPU 2		CPU 3	
CORE 2	CPU 4		CPU 5	
CORE 3	CPU 6		CPU 7	
CORE 4	CPU 8		CPU 9	
CORE 5	CPU 10		CPU 11	

```
int numMaxCores = sysInfo.wProcessorLevel;    // 6 in our case
int numMaxThreads = sysInfo.dwNumberOfProcessors; // 12 in our case

if (NR_THREADS <= numMaxThreads) {
    // Set CPU affinity for each thread
    for (int i = 0; i < NR_THREADS; i++) {
        // Assign each thread to a different CPU logical core
        DWORD_PTR mask = 1 << i;
        if (!SetThreadAffinityMask(handles[i], mask)) {
            fprintf(stderr, "Error setting affinity for thread %d\n", i);
            return 1;
        }
    }
}
```

# Following logical cores order



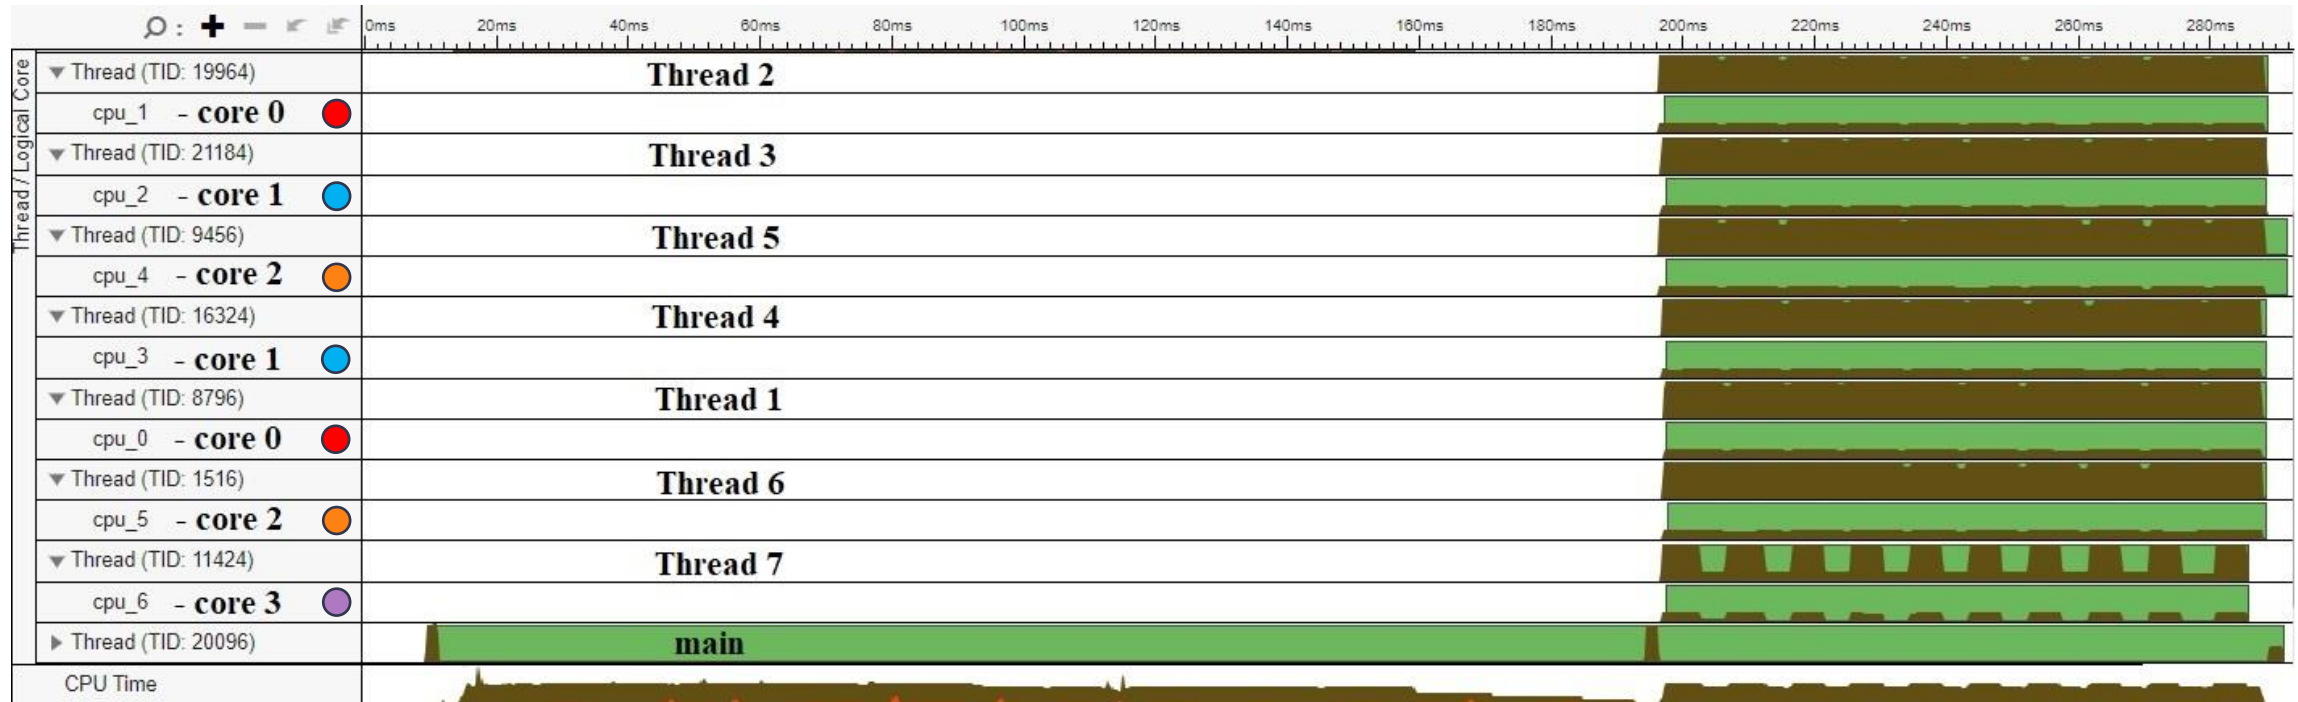
- 100 000 Points dataset

- 7 Threads

- Clusters=5 & Max # of iterations = 10

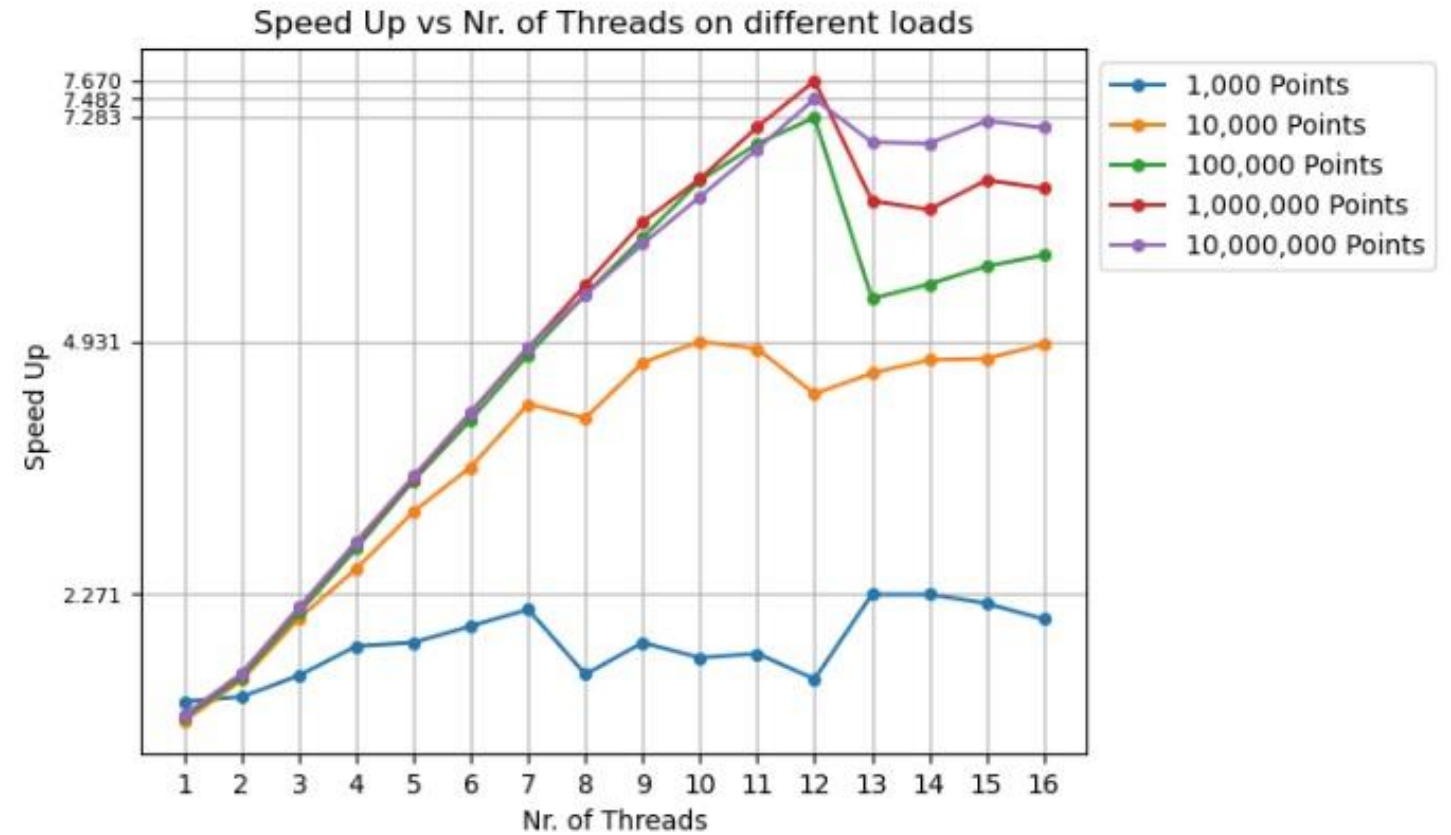
● Running

● CPU in use



# Speed Up





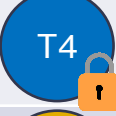
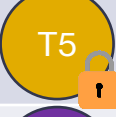
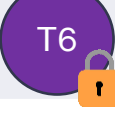
affinity set to follow  
logical cores order



# Using affinity



# Following physical cores order

CORE 0	CPU 0		CPU 1	
CORE 1	CPU 2		CPU 3	
CORE 2	CPU 4		CPU 5	
CORE 3	CPU 6		CPU 7	
CORE 4	CPU 8		CPU 9	
CORE 5	CPU 10		CPU 11	

```
int numMaxCores = sysInfo.wProcessorLevel;    // 6 in our case
int numMaxThreads = sysInfo.dwNumberOfProcessors; // 12 in our case

if (NR_THREADS <= numMaxThreads) {
    // Set CPU affinity for each thread
    for (int i = 0; i < NR_THREADS; i++) {

        // Assign each thread to a different CPU logical core

        int shift = ((i * 2) < numMaxThreads) ? (i * 2) : ((i % numMaxCores) * 2 + 1);

        DWORD_PTR mask = 1 << shift;
        if (!SetThreadAffinityMask(handles[i], mask)) {
            fprintf(stderr, "Error setting affinity for thread %d\n", i);
            return 1;
        }
    }
}
```

# Following physical cores order

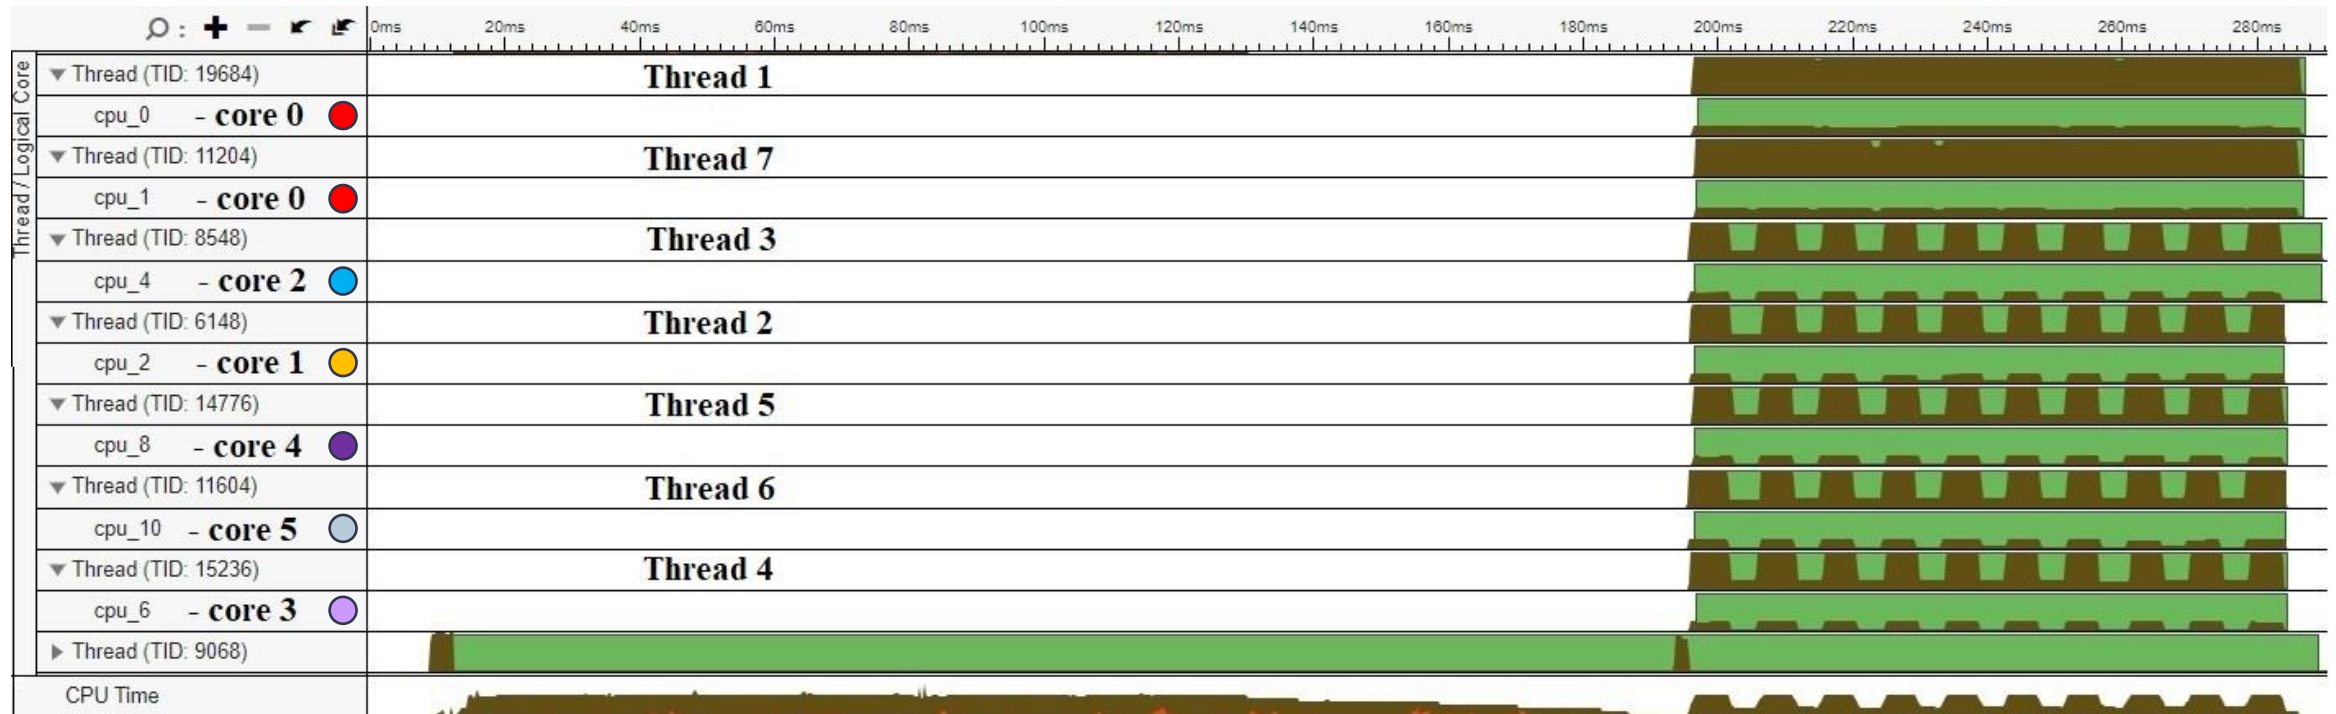


- 100 000 Points dataset

- 7 Threads

- Clusters=5 & Max # of iterations = 10

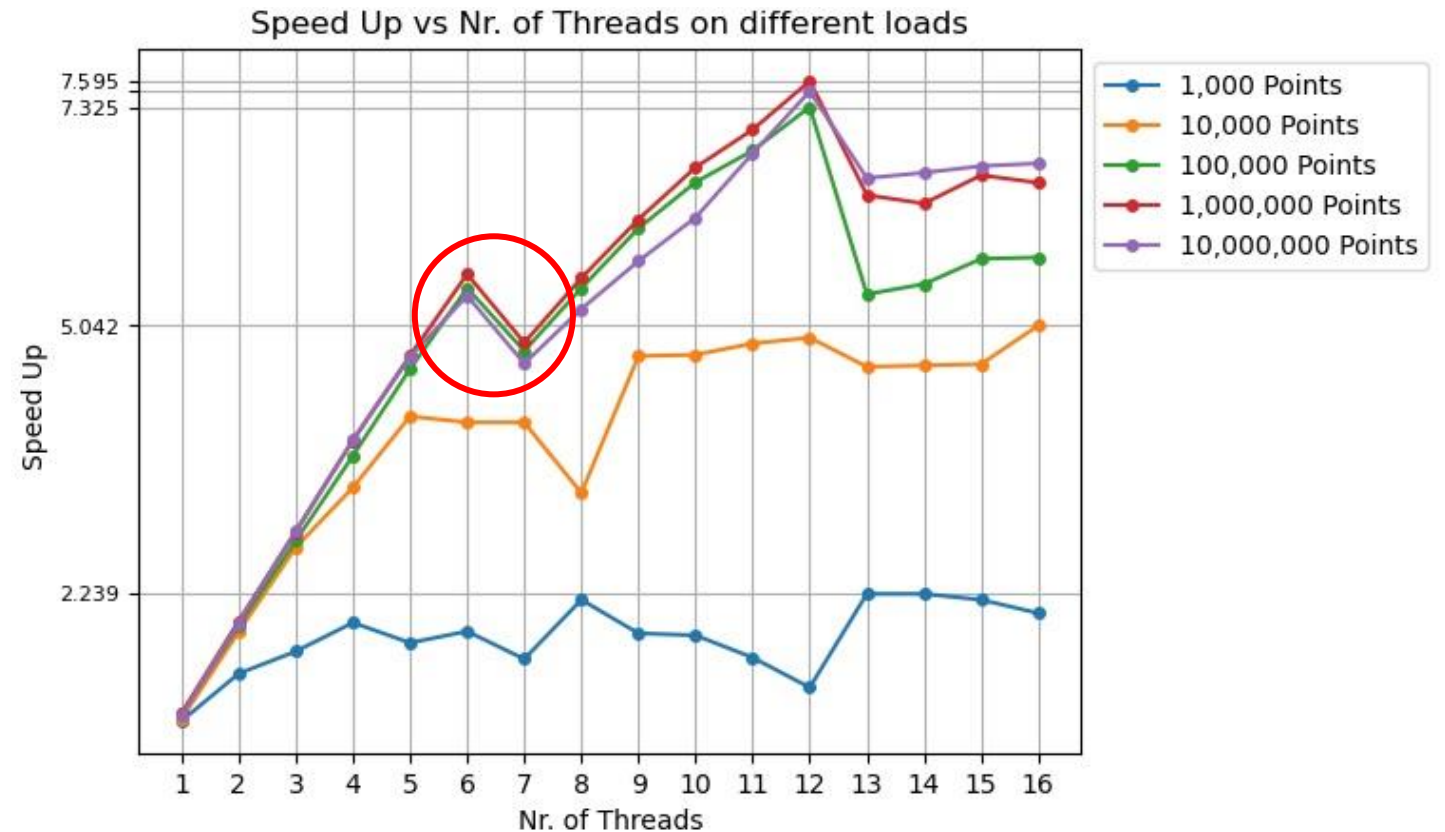
● Running      ● CPU in use





# Speed Up

affinity set to follow  
physical cores order





# What's the cause of this effect?

**Memory bound** : issues on the memory subsystem (cache) that causes the introduction of stalls inside pipeline slots due miss.

# How to interpret the values



Physical Core / Thread / Function / Call Stack	CPU Time ▼	Memory Bound					LLC Miss Count	Loads	Stores
		L1 Bound	L2 Bound	L3 Bound	DRAM Bound	Store Bound			
core_3	134.242ms	22.1%	0.9%	0.5%	0.5%	0.1%	26,000	93,847,000	31,655,000
core_0	101.974ms	20.0%	0.4%	2.0%	0.0%	0.1%	0	68,471,000	23,465,000
core_5	95.614ms	20.5%	0.6%	1.3%	0.0%	0.0%	0	68,627,000	23,218,000
core_1	95.175ms	19.1%	0.7%	1.1%	0.0%	0.0%	0	75,413,000	25,441,000
core_4	83.709ms	19.8%	0.7%	0.7%	0.4%	0.1%	0	62,803,000	21,190,000
core_2	74.812ms	17.9%	0.6%	0.5%	0.4%	0.2%	0	57,551,000	19,422,000

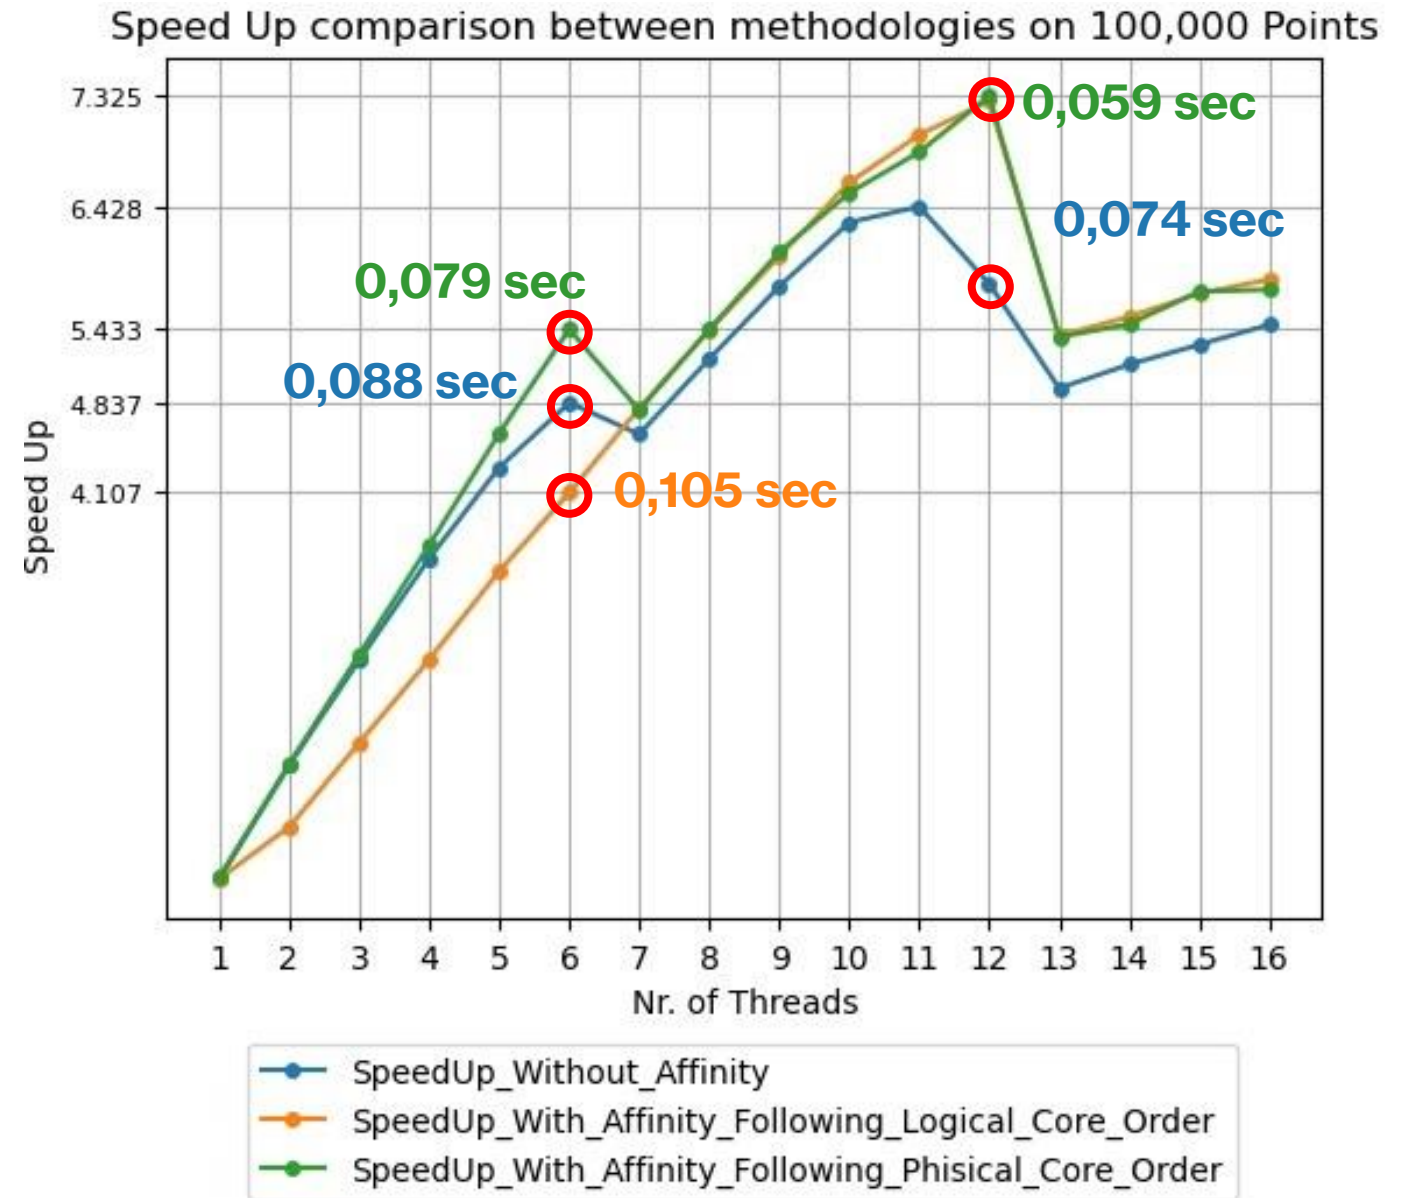
- 100 000 Points dataset

- 6 Threads

- Without Affinity

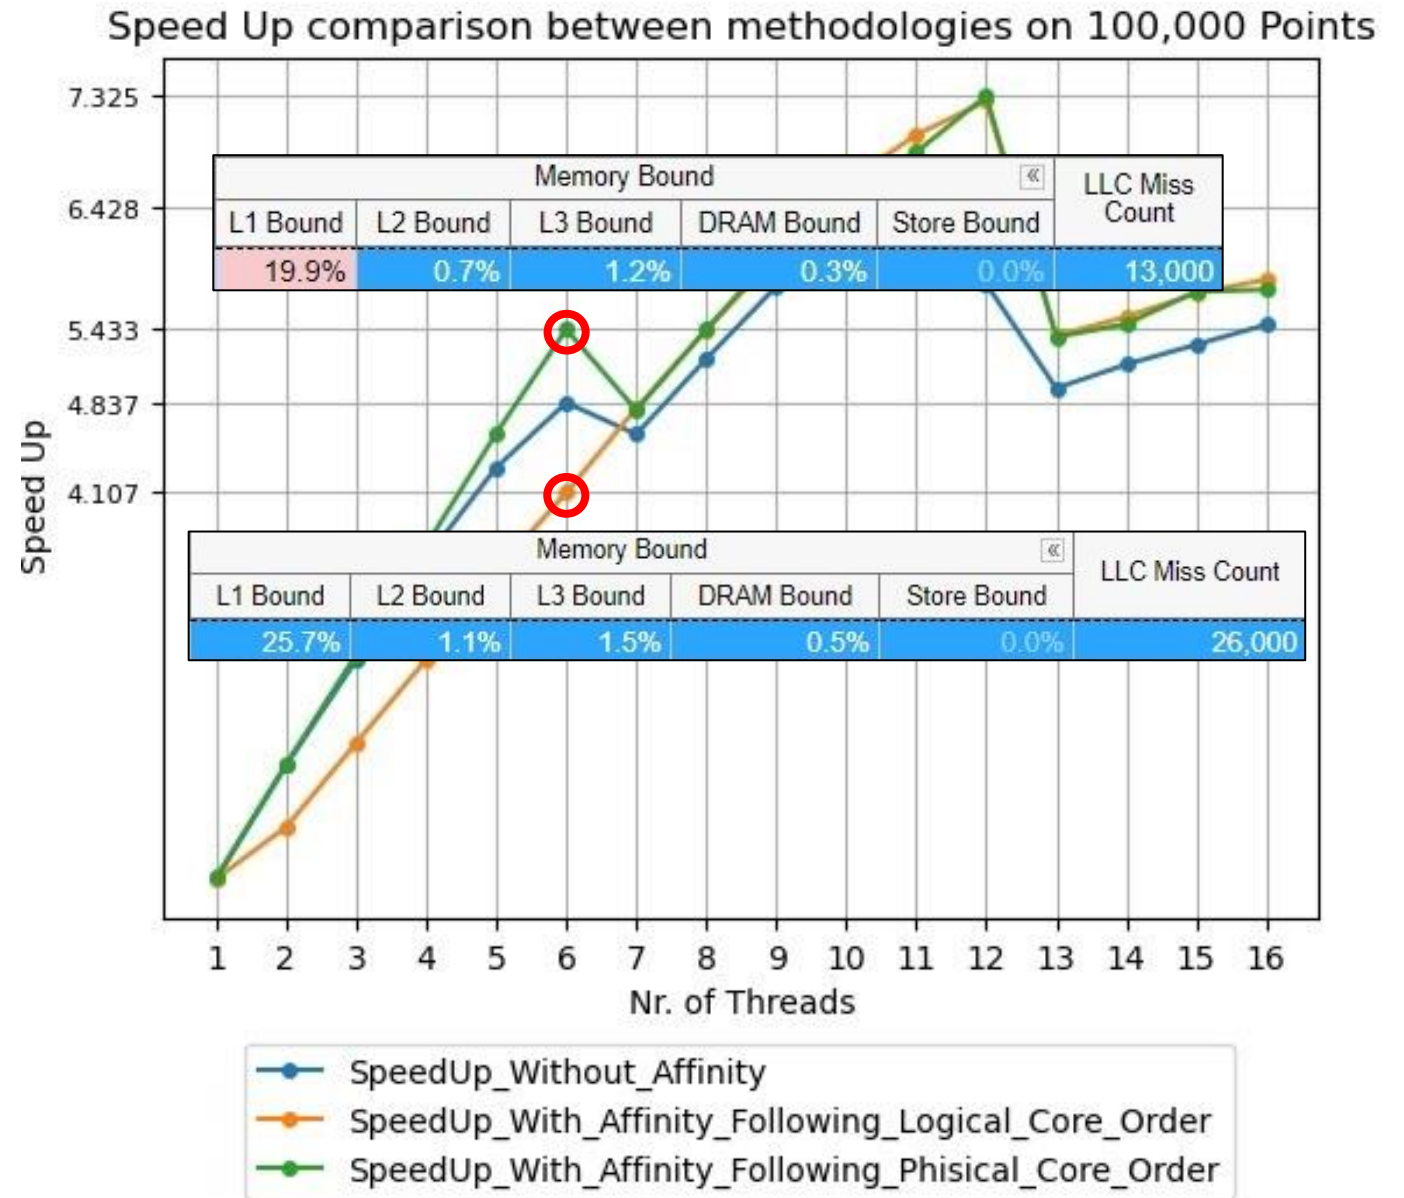
# Speed Up

Comparison on  
100.000 points with  
different affinity  
methodologies



# Speed Up

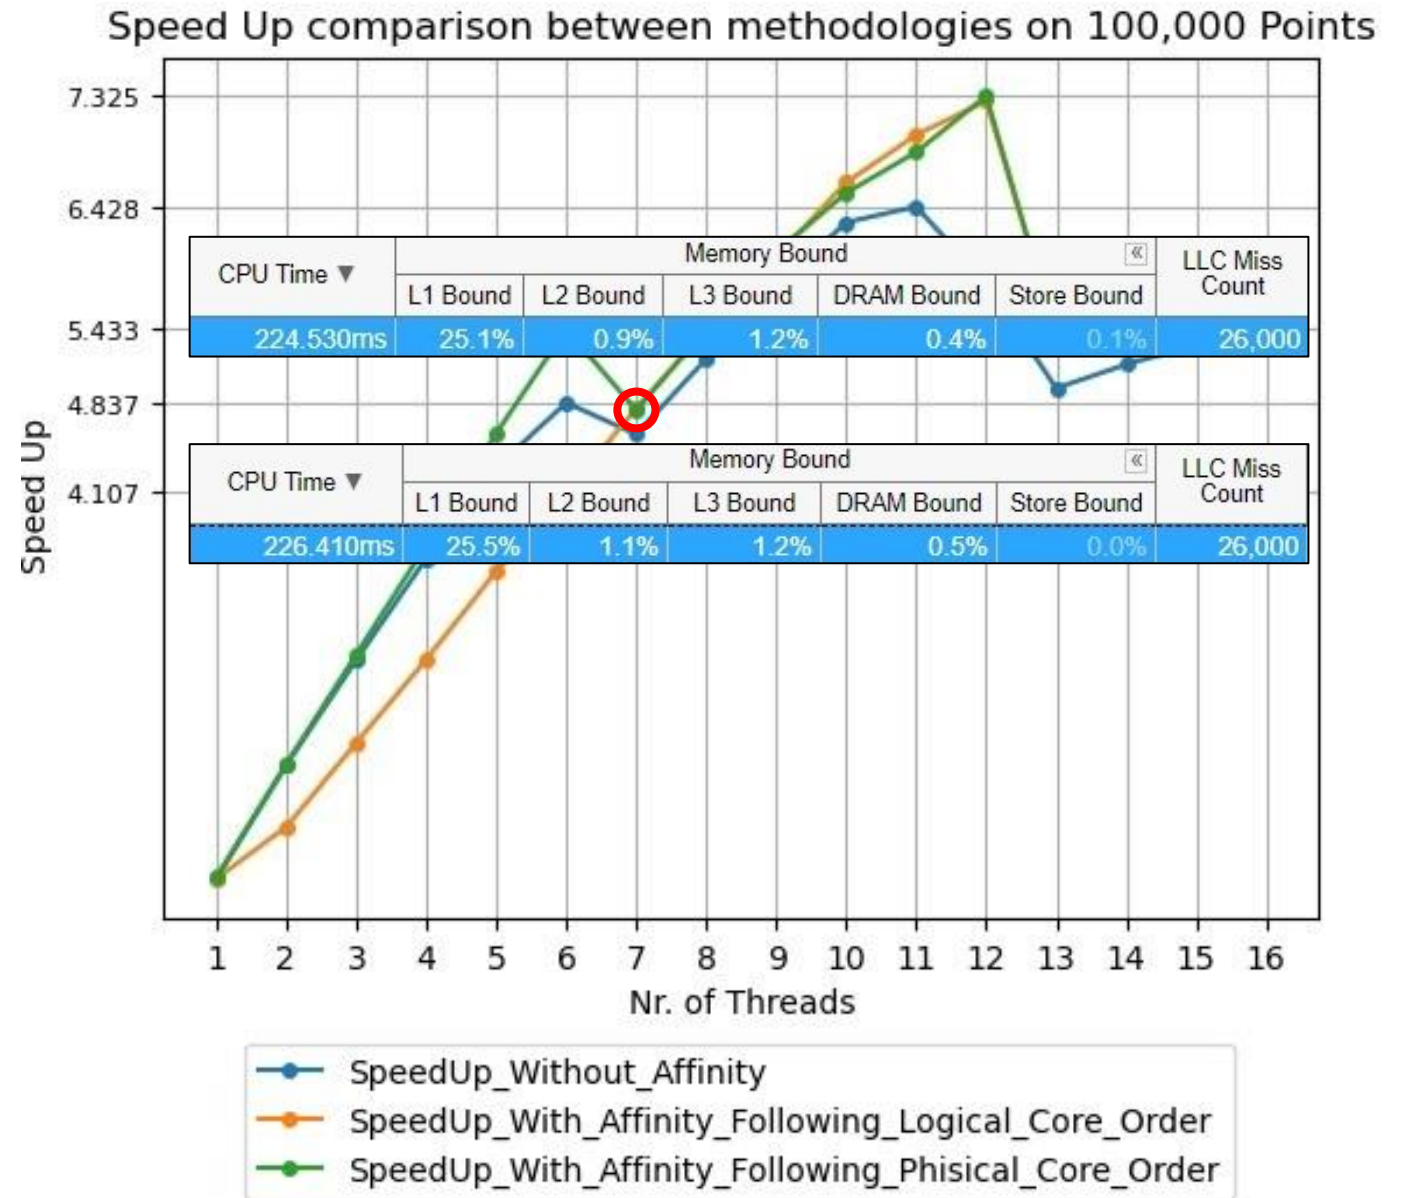
Comparison on  
100.000 points with  
different affinity  
methodologies





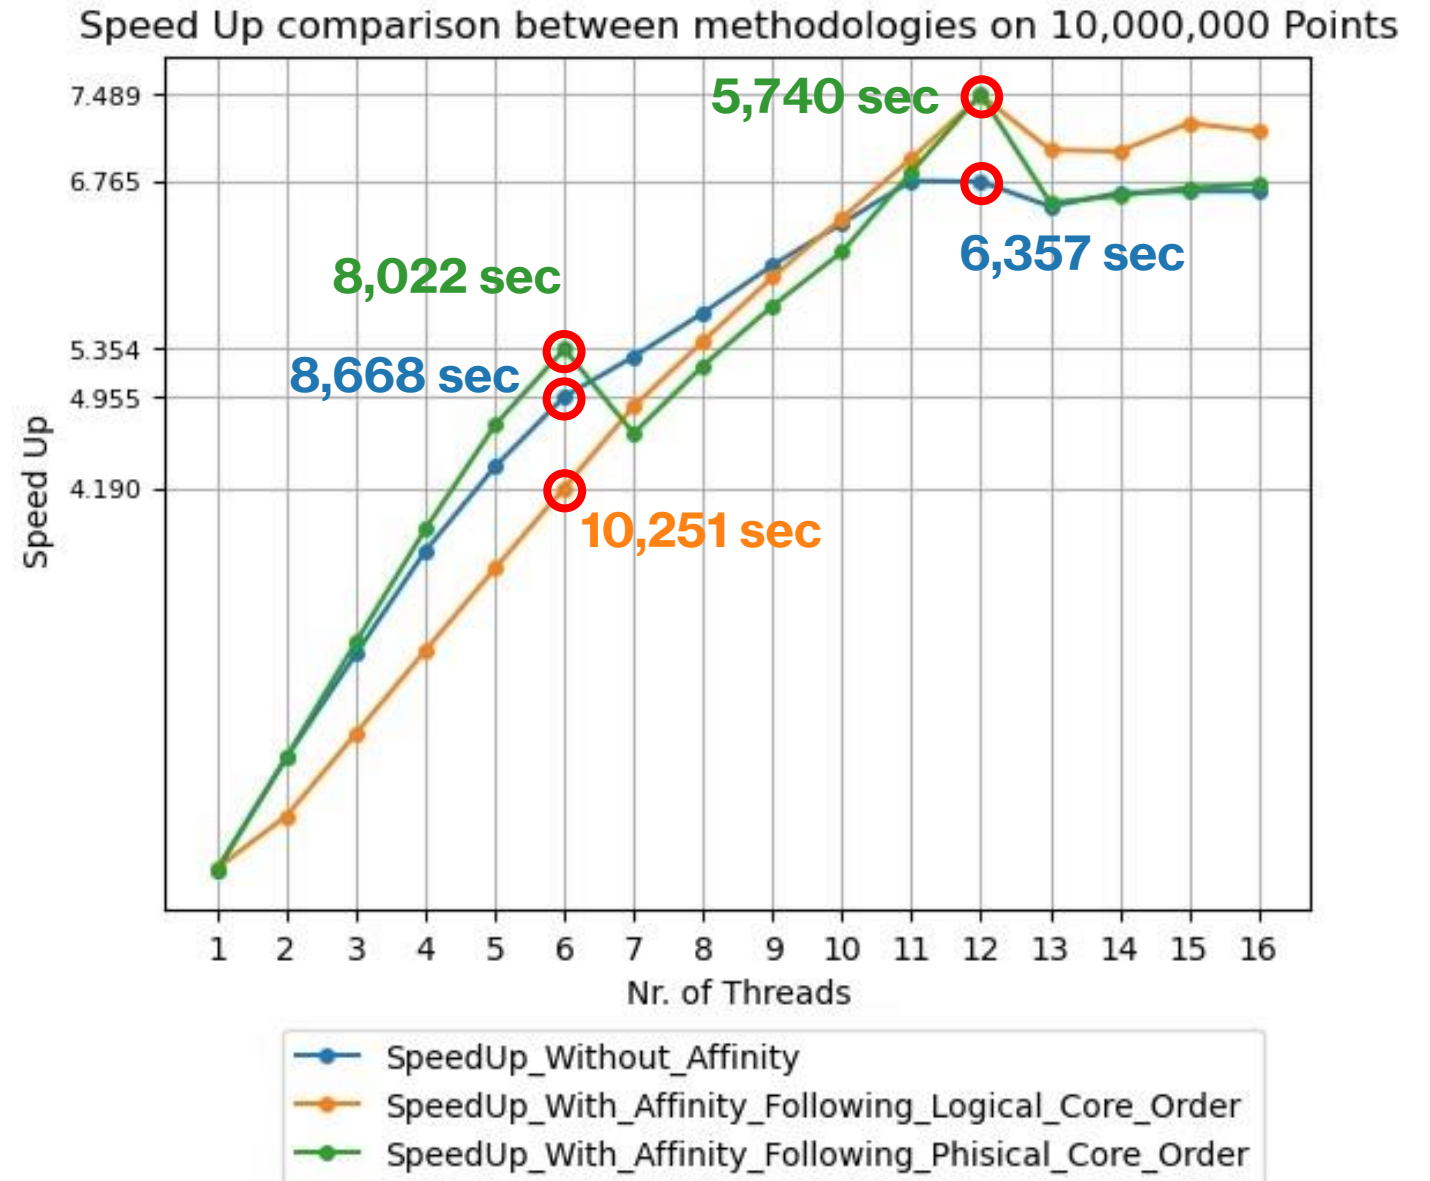
# Speed Up

Comparison on  
100.000 points with  
different affinity  
methodologies



# Speed Up

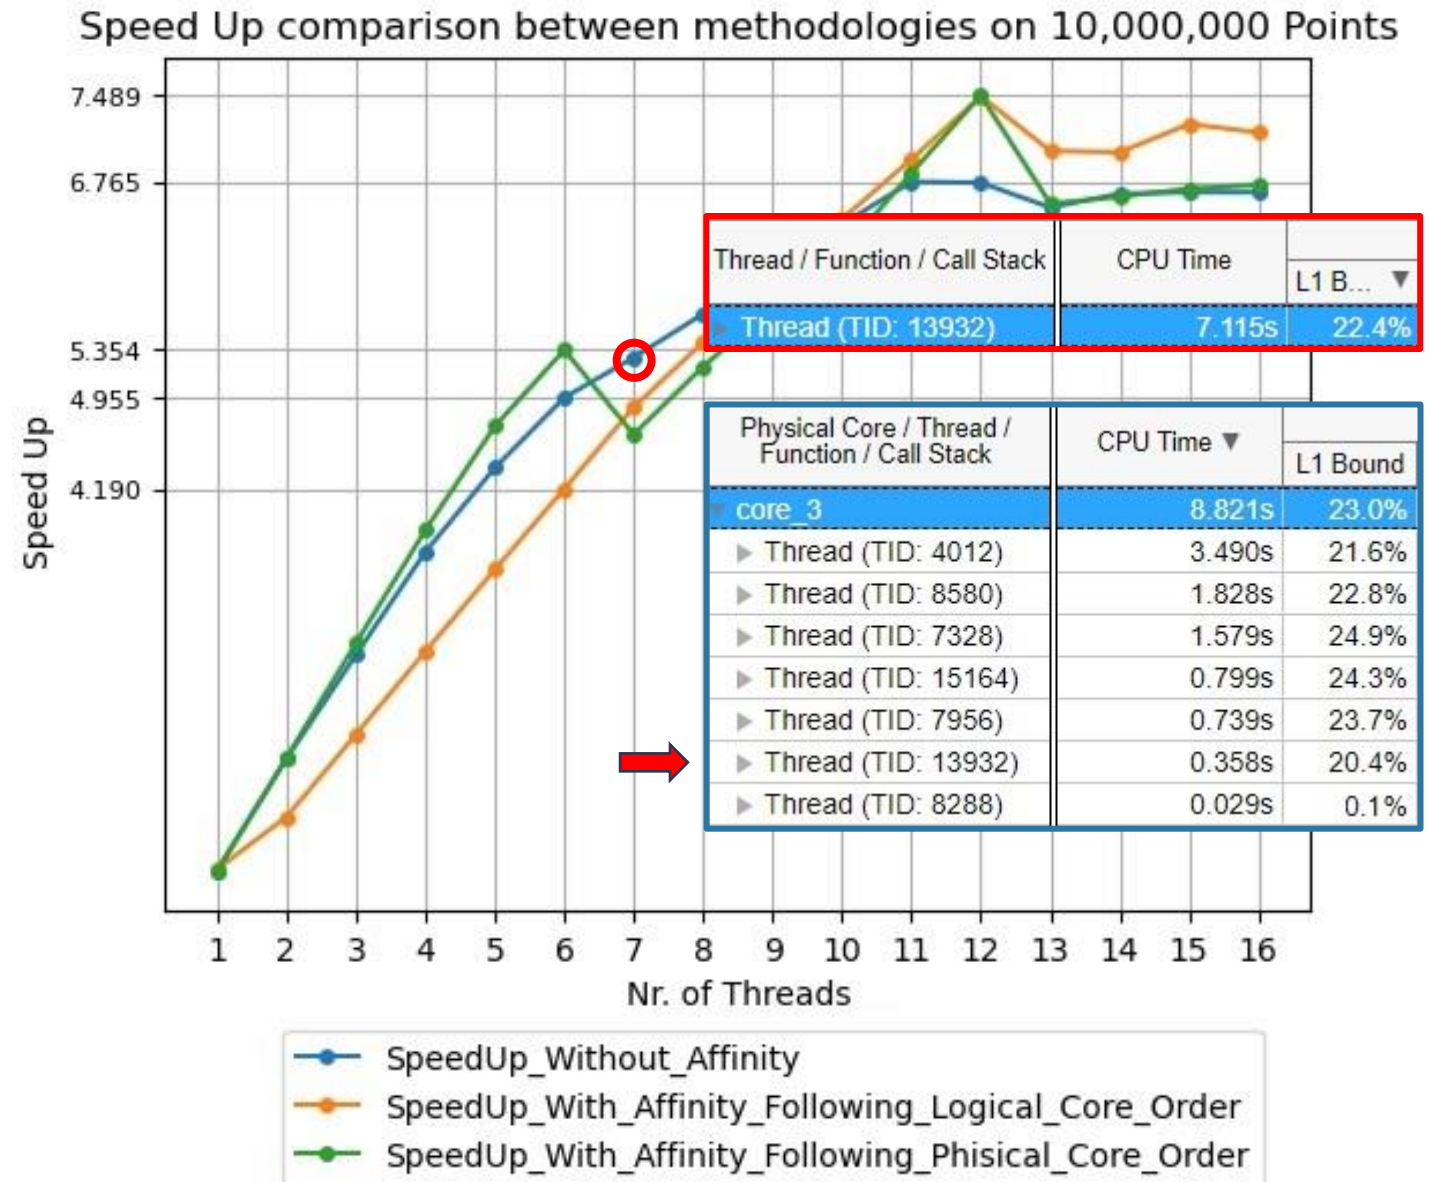
Comparison on  
10.000.000 points  
with different affinity  
methodologies



# Speed Up

## Comparison on 10.000.000 points with different affinity methodologies

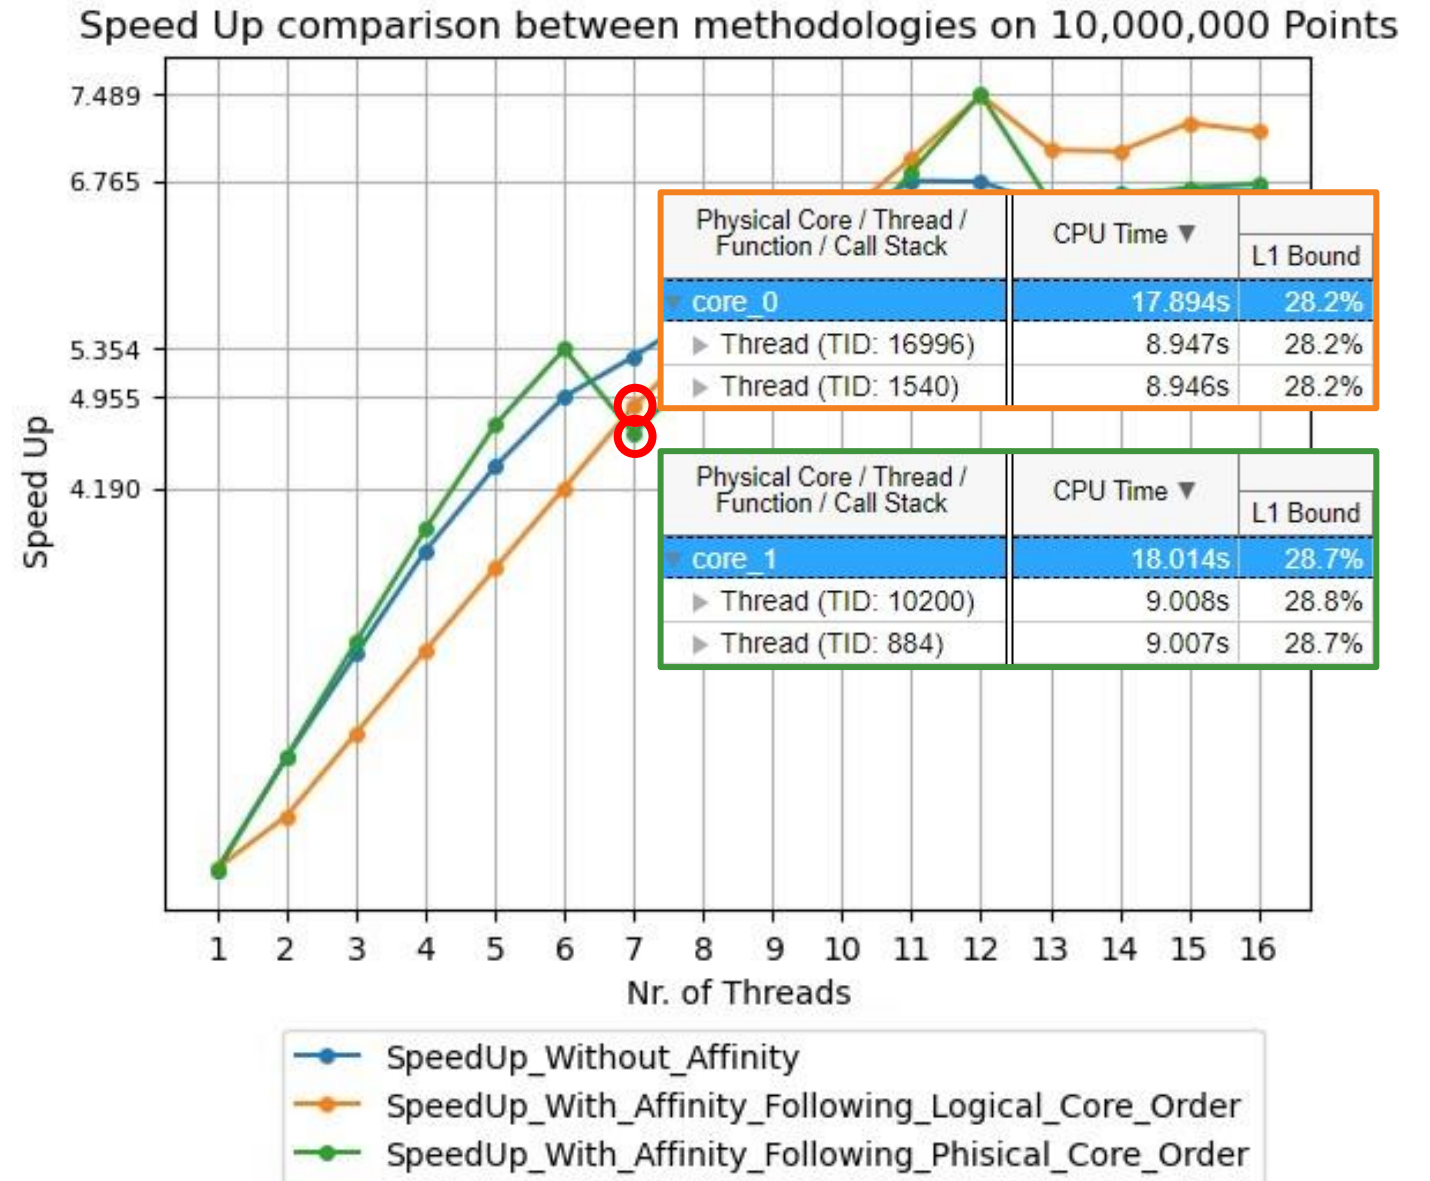
The O.S obtains a smoother curve increasing the load and in some cases better than with fixed affinity management.



# Speed Up

## Comparison on 10.000.000 points with different affinity methodologies

The O.S obtains a smoother curve increasing the load and in some cases better than with fixed affinity management.





# Optimizations

## OPT.1: Data structures

```
// Define the structure for a Point
typedef struct {
    double coordinate[3];
} Point;
```

```
DynamicList points;
DynamicList centroids;
DynamicList sums;
```



```
// Define the structure for the dynamic list
typedef struct {
    Point* points;
    size_t capacity;
    size_t size;
} DynamicList;

// METHODS TO HANDLE DynamicList TYPE
void initDynamicList(DynamicList* list, size_t initial_capacity);
void appendPoint(DynamicList* list, Point point);
Point* accessByIndex(DynamicList* list, size_t index);
void freeDynamicList(DynamicList* list);
```

```
Point* points = (Point*)malloc(DATASET_SIZE * sizeof(Point));
Point* centroids = (Point*)malloc(K * sizeof(Point));
Point* sums = (Point*)malloc(K * sizeof(Point));
```

# Optimizations

## OPT.2: Math.h

```
for (int clusterIndex = 0; clusterIndex < arguments->K; clusterIndex++) {  
    float sumPartial = 0;  
    for (int dim = 0; dim < arguments->DIM; dim++) {  
        sumPartial += pow(arguments->centroids[clusterIndex].coordinate[dim] - point.coordinate[dim], 2);  
    }  
    // Compute distance from a point to all centroids  
    dists[clusterIndex] = sqrt(sumPartial);  
}
```



```
for (int clusterIndex = 0; clusterIndex < arguments->K; clusterIndex++) {  
    float sumPartial = 0;  
    for (int dim = 0; dim < arguments->DIM; dim++) {  
        double partialDiff = arguments->centroids[clusterIndex].coordinate[dim] - point.coordinate[dim];  
        sumPartial += partialDiff * partialDiff;  
    }  
    // Compute distance from a point to all centroids  
    dists[clusterIndex] = sqrt(sumPartial);  
}
```



# Optimizations

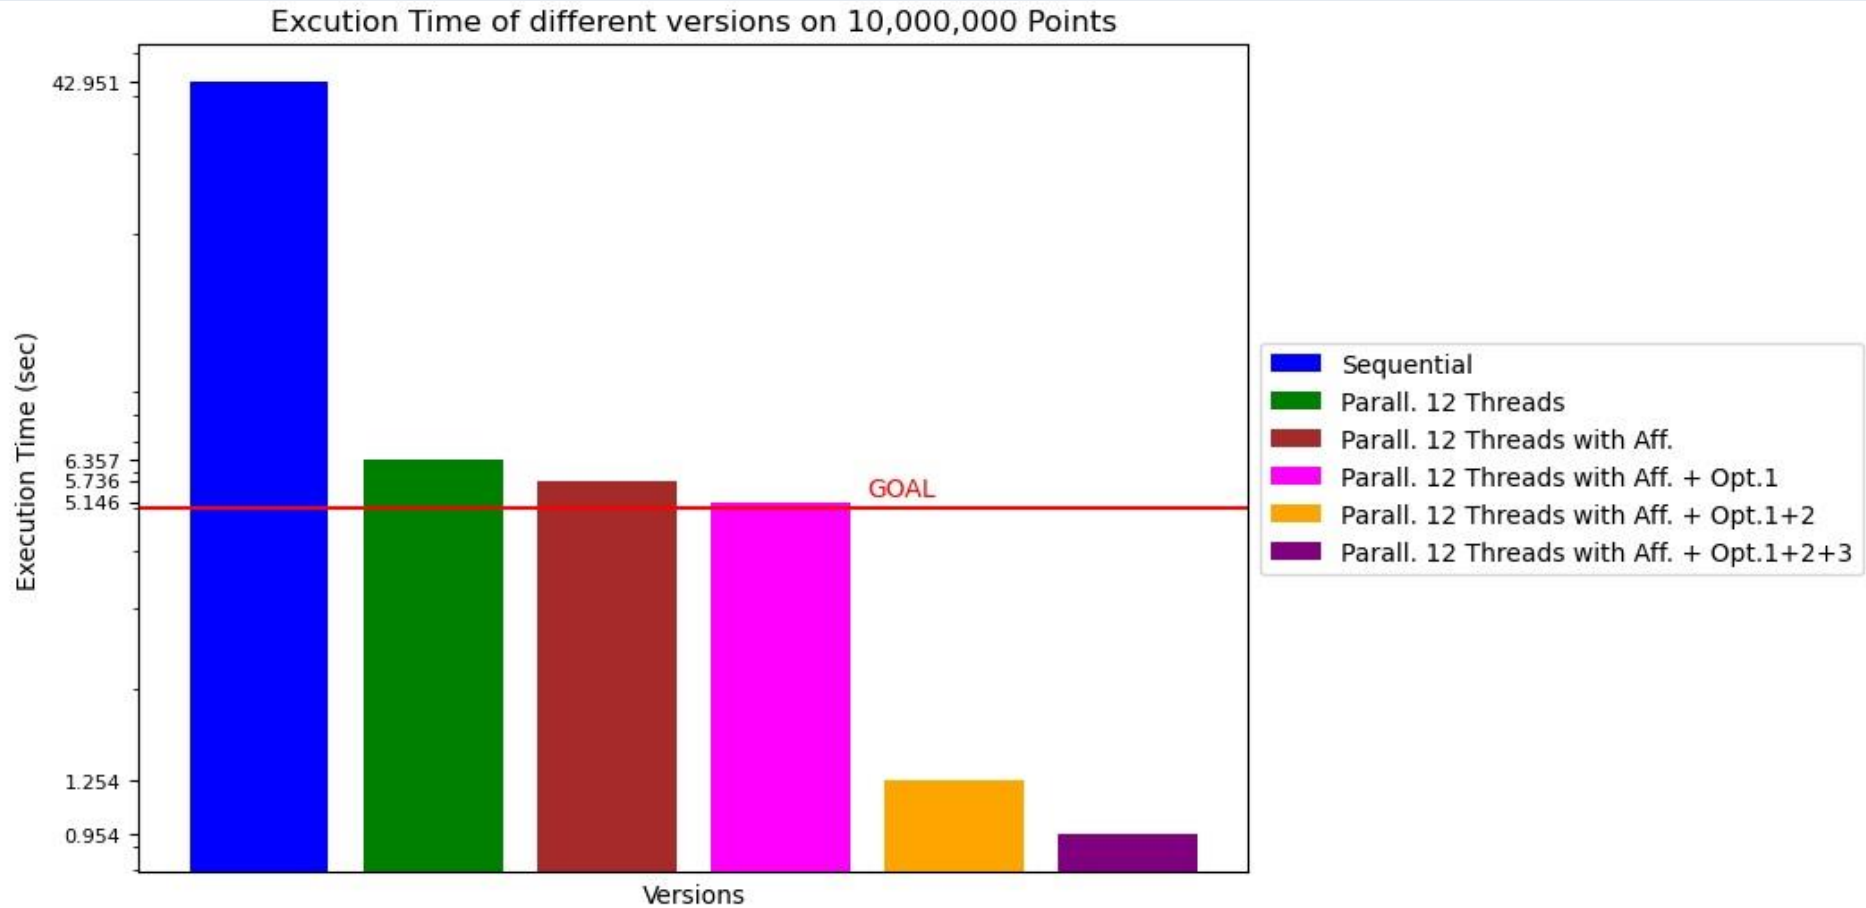
## OPT.3: To act on the algorithm ... a little

We are not actually interested in obtaining the exact distance, but only to discover the minimum among all the centroids.

We can avoid to compute the square root.

```
for (int clusterIndex = 0; clusterIndex < arguments->K; clusterIndex++) {  
    float sumPartial = 0;  
    for (int dim = 0; dim < arguments->DIM; dim++) {  
        double partialDiff = arguments->centroids[clusterIndex].coordinate[dim] - point.coordinate[dim];  
        sumPartial += partialDiff * partialDiff;  
    }  
    // Compute distance from a point to all centroids  
    /*dists[clusterIndex] = sqrt(sumPartial);*/  
    dists[clusterIndex] = sumPartial;  
}
```

# Optimizations results





# If we use a GPU?



# GPU Nvidia GeForce GTX 1060



```
Device 0: "NVIDIA GeForce GTX 1060 6GB"
  CUDA Driver Version / Runtime Version      12.4 / 12.4
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             6144 MBytes (6442188800 bytes)
  (010) Multiprocessors, (128) CUDA Cores/MP: 1280 CUDA Cores
  GPU Max Clock rate:                        1759 MHz (1.76 GHz)
  Memory Clock rate:                         4004 Mhz
  Memory Bus Width:                          192-bit
  L2 Cache Size:                             1572864 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total shared memory per multiprocessor:     98304 bytes
  Total number of registers available per block: 65536
  Warp size:                                 32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                          512 bytes
```

# GPU – Initial Version

We maintain the same exact structure of the parallel CPU version

- 2 kernel functions

```
1) assignPointsToClusterKernel<< <blocks, threads >> > (d_points, d_centroids,  
d_membership, d_sums, d_counts, SIZE);
```

```
2) updateCentroidsKernel << <1, K >> > (d_centroids, d_sums, d_counts);
```

- Same measurement methods: function clock() in time.h C library
  - Focus on pure execution time
- What is its behavior? How do we select the # of blocks and threads?

# GPU – Initial Version

```
printf("\nStart of the algorithm :");

start_alg = clock();

// Main loop of KMeans algorithm
for (int iter = 0; iter < MAX_ITERATIONS; ++iter) {

    // Launch kernel to assign points to clusters
    assignPointsToClusterKernel << <blocks, threads >> > (d_points, d_centroids, d_membership, d_sums, d_counts, SIZE);
    cudaDeviceSynchronize();

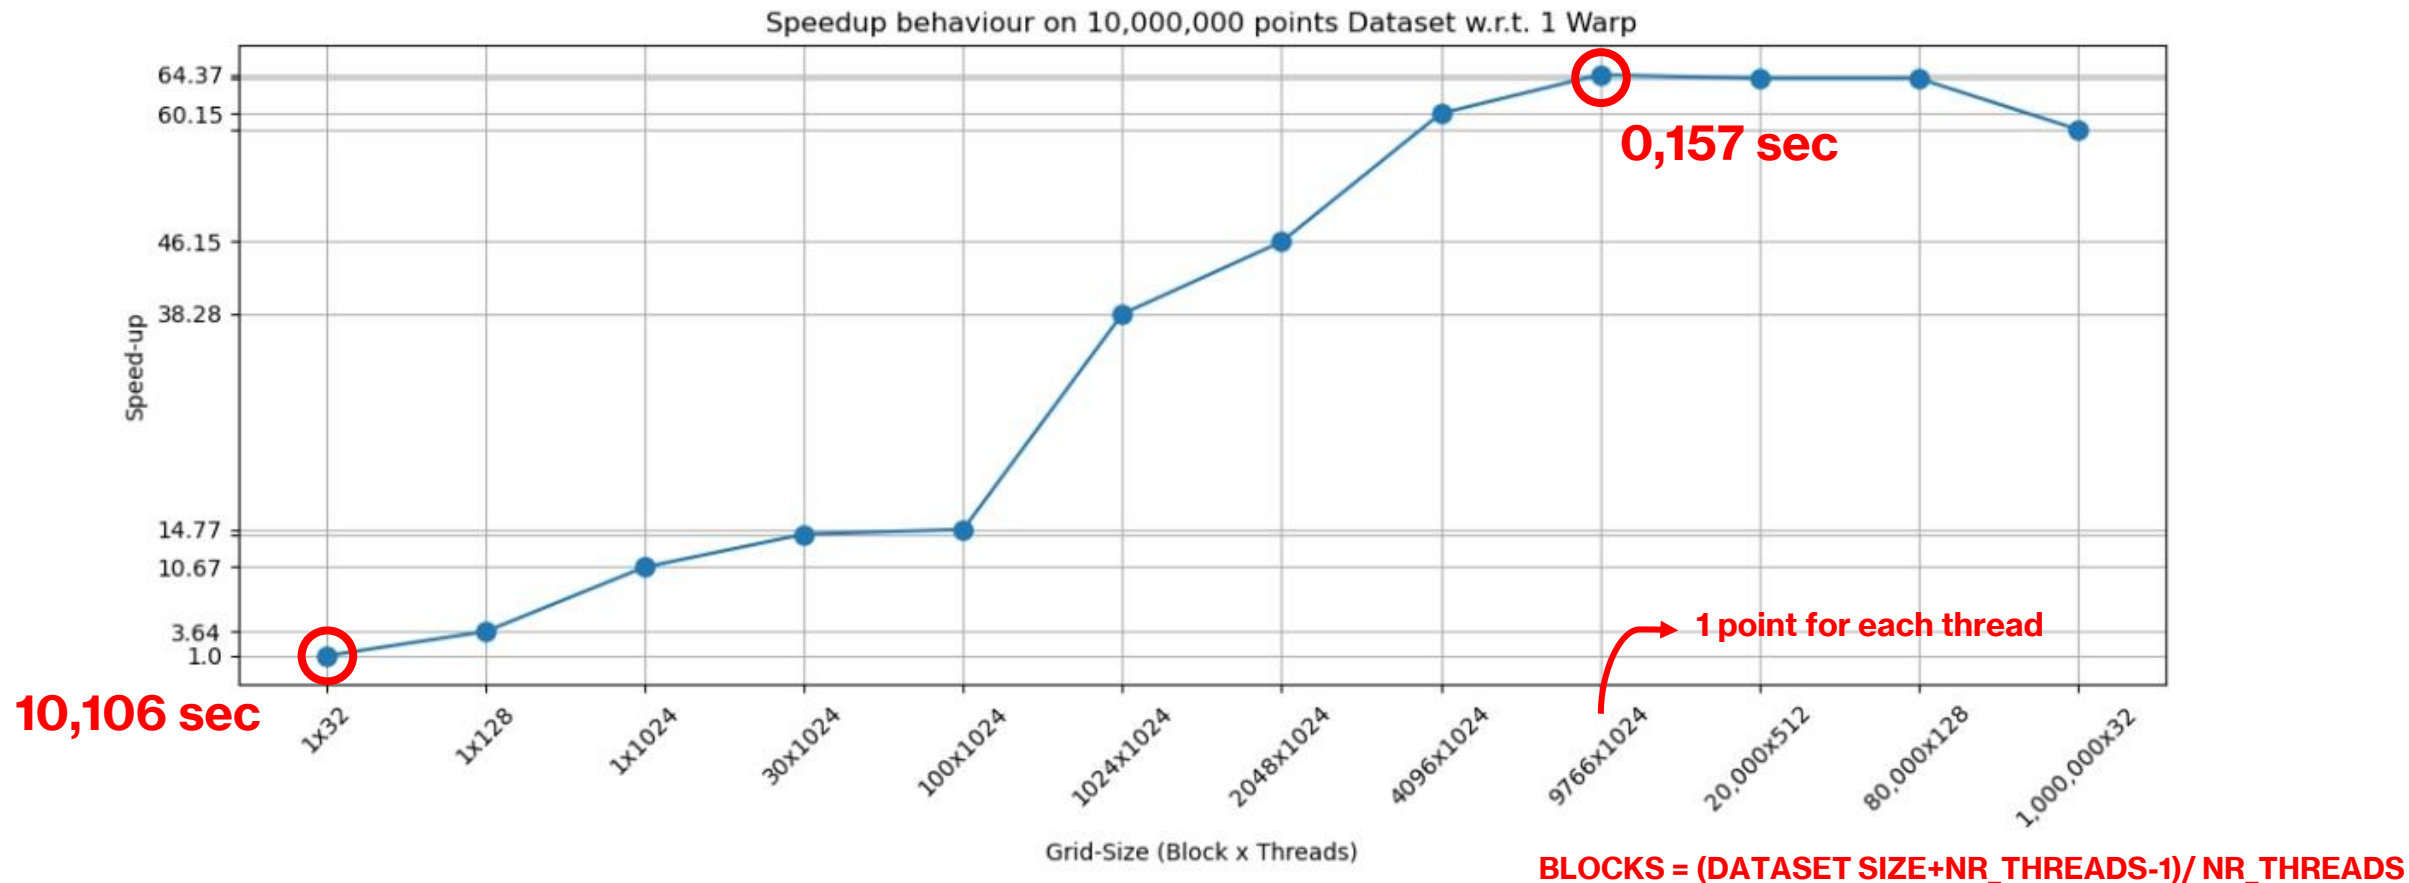
    // Launch kernel to update centroids
    updateCentroidsKernel << <1, K >> > (d_centroids, d_sums, d_counts);
    cudaDeviceSynchronize();
}

end_alg = clock();

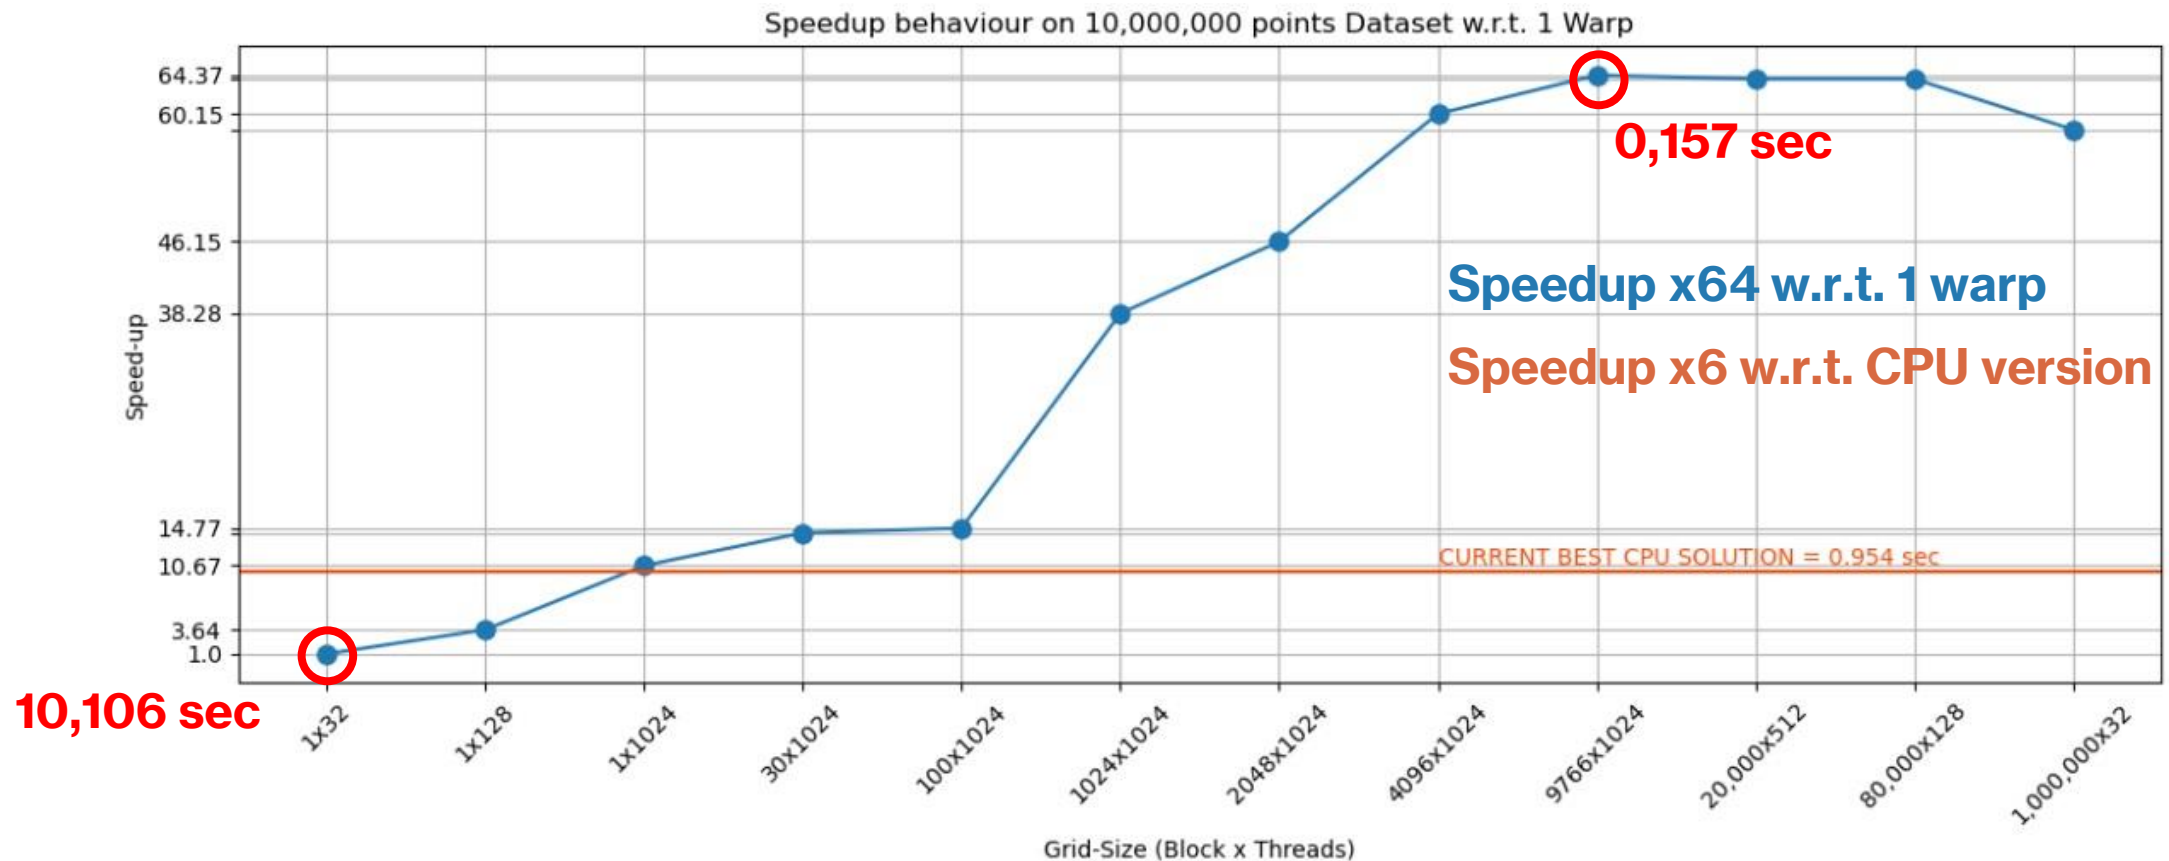
printf("End of the algorithm: \n");
```



# GPU – Initial Version



# GPU – Initial Version



# GPU – Profiler



NVIDIA  
Nsight™  
Compute

v.2019

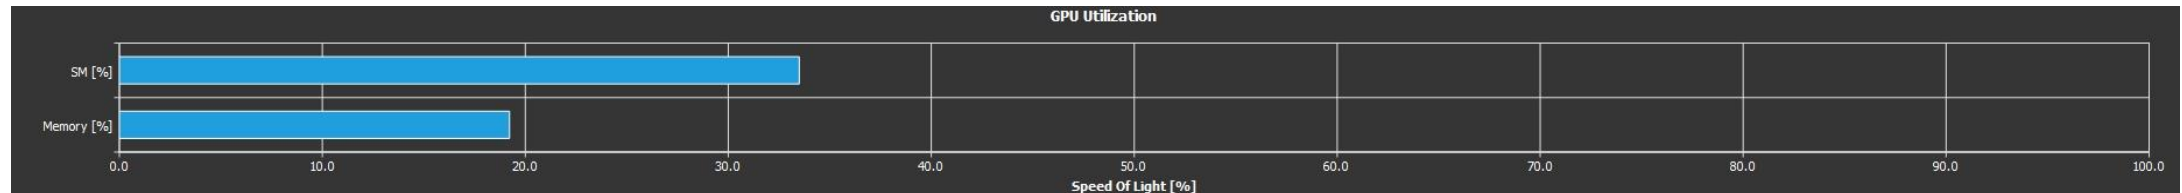
## Recommendations



### Bottleneck

[Warning] This kernel exhibits low compute throughput and memory bandwidth utilization relative to the peak performance of this device. Achieved compute throughput and/or memory bandwidth below 60.0% of peak typically indicate latency issues. Look at "Scheduler Statistics" and "Warp State Statistics" for potential reasons.

SOL SM [%]	33,50
SOL Memory [%]	19,22
SOL TEX [%]	15,60
SOL L2 [%]	19,22
SOL FB [%]	8,73



# GPU – Profiler



NVIDIA  
Nsight™  
Compute

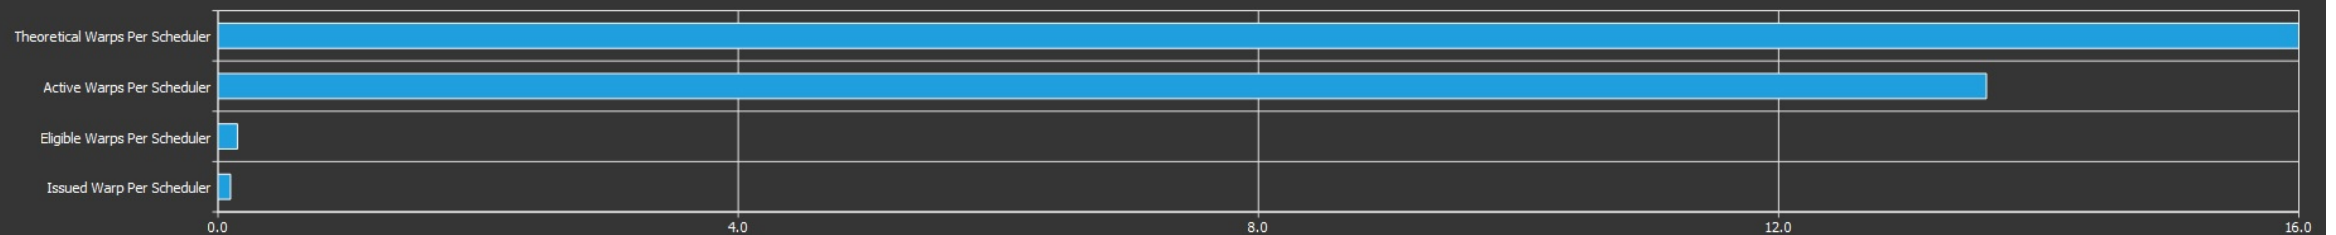
v.2019

## Recommendations

**⚠ Issue Slot Utilization** [Warning] Every scheduler is capable of issuing two instructions per cycle, but for this kernel each scheduler only issues an instruction every 10.4 cycles. This might leave hardware resources underutilized and may lead to less optimal performance. Out of the maximum of 16 warps per scheduler, this kernel allocates an average of 13.60 active warps per scheduler, but only an average of 0.15 warps were eligible per cycle. Eligible warps are the subset of active warps that are ready to issue their next instruction. Every cycle with no eligible warp results in no instruction being issued and the issue slot remains unused. To increase the number of eligible warps either increase the number of active warps or reduce the time the active warps are stalled.

Active Warps Per Scheduler [warp/cycle]	13,60
Eligible Warps Per Scheduler [warp/cycle]	0,15
Issued Warp Per Scheduler [issue/cycle]	0,10
Instructions Per Active Issue Slot [inst/issue]	1,07
No Eligible [%]	90,26
One or More Eligible [%]	9,63

Warps Per Scheduler



# GPU – Considerations

Data structure used:

```
typedef struct {  
    double coordinate[3];  
} Point;  
Point* points = (Point*)malloc(DATASET_SIZE_HOST * sizeof(Point));
```

Operation performed:

```
// Update sums and counts  
for (int dim = 0; dim < DIM; dim++) {  
    atomicAdd(&sums[minIndex].coordinate[dim], points[indexPartial].coordinate[dim]);  
}
```

Atomic operations on double generate errors at compilation time!!

How to solve?

Two options:

1. Hardware option: `nvcc kernel.cu -o KMeansCuda -arch=sm_61` (Faster)
2. Software (Slower)

**We used this one!**



<https://docs.nvidia.com/cuda/cuda-compiler-driver-nvcc/index.html#gpu-feature-list>

# GPU – Considerations

## Memory usage

- Constant memory
- Shared Memory

```
// Algorithm parameters used inside the device.
__constant__ int K;
__constant__ int DIM;
__constant__ int DATASET_SIZE;

// Set constant values
cudaMemcpyToSymbol(K, &K_HOST, sizeof(int));
cudaMemcpyToSymbol(DIM, &DIM_HOST, sizeof(int));
cudaMemcpyToSymbol(DATASET_SIZE, &DATASET_SIZE_HOST, sizeof(int));
```

```
// Inside main function
// Launch kernel to assign points to clusters
assignPointsToClusterKernel << <blocks, threads , K_HOST *sizeof(Point) >> >
cudaDeviceSynchronize();

// Inside kernel function
// Dynamic shared memory for centroids among blocks
extern __shared__ Point sharedCentroids[];

// Load centroids into shared memory for the first k and wait the others
if (threadIdx.x < K) {
    sharedCentroids[threadIdx.x] = centroids[threadIdx.x];
}
__syncthreads();
```



# GPU – Considerations

To maximize the usage of the architecture, considering that we are operating with bi-dimensional points, let's exploit gpu features, and use a 2D grid, in order to do so, we need to change the data structure previously used, to a 2D matrix.

**2D Matrix  
of doubles**

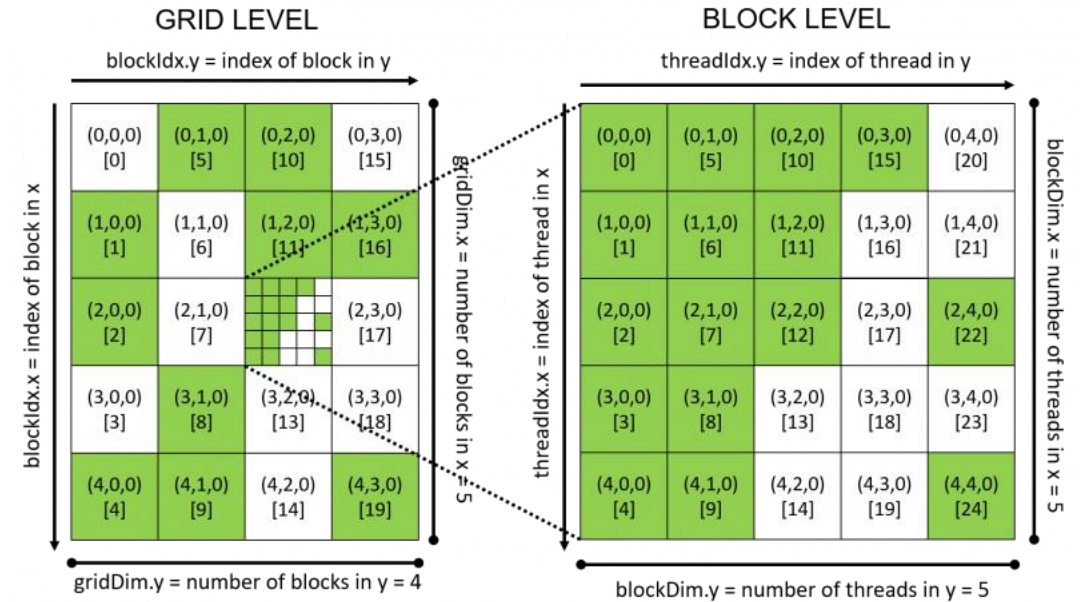


```
double** pointsMatrix;  
pointsMatrix = (double**)malloc(sizeof(double**) * DIM);  
for (int i = 0; i < DIM; i++)  
    pointsMatrix[i] = (double*)malloc((sizeof(double*) * DATASET_SIZE));
```

# GPU – Considerations

## 2D-GRID

```
double** d_pointsMatrix;  
double** d_pointsR;  
d_pointsR = (double**)malloc(DIM * sizeof(double*));  
  
cudaMalloc((void**)&d_pointsMatrix, DIM * sizeof(double*));  
  
for (int i = 0; i < DIM; i++)  
    cudaMalloc((void**)&d_pointsR[i], DATASET_SIZE * sizeof(double));
```



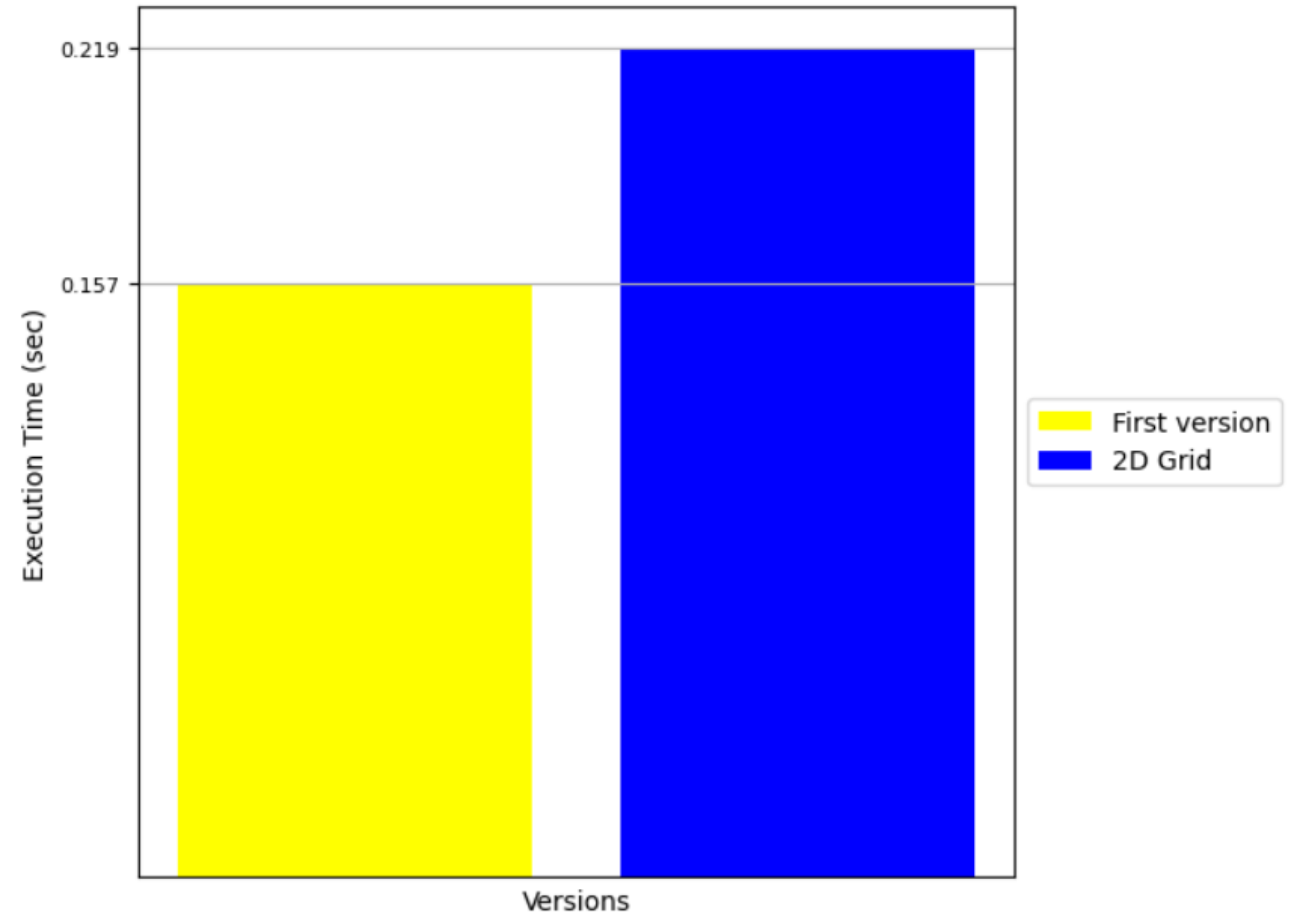
# GPU – Considerations

We used block size large enough to cover the entire dataset, and in order to provide enough warps to schedule, using the architecture at it's maximum.

```
//with this values it's possible to modify blocks size
int dimBlockX = 32;
int dimBlockY = 32;
// this values are calculated in order to use the minimum number of threads to cover the entire dataset
// and to use matrix of N X N
dim3 blockSize(dimBlockX, dimBlockY);
int dimGrid = ceil(sqrt(((DATASET_SIZE - (dimBlockX * dimBlockY) - 1)) / (dimBlockX * dimBlockY)));
int dimGridX = dimGrid;
int dimGridY = dimGrid;
dim3 gridSize(dimGridX, dimGridY);

assignPointsToClusterKernel << ~ gridSize, blockSize >> >
```

# Execution time comparison between Initial version and using a 2D Grid



# GPU – Profiler

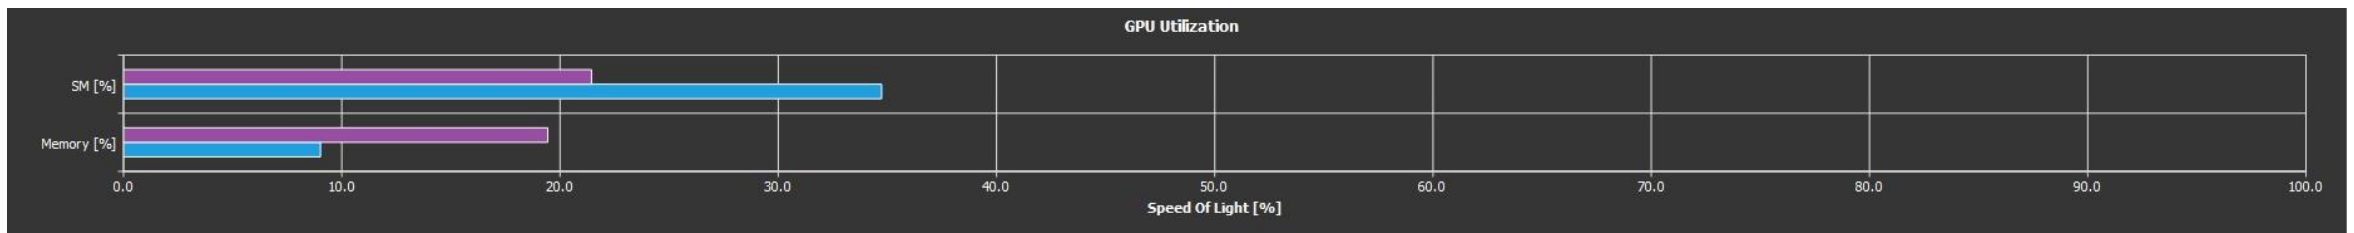


v.2019

We reached the goal of increasing the usage of the Multiprocessors (+60% SM utilization) but the exploiting of the memory subsystem of the cache is worse with the new version (-53%)

● Initial Version

● 2D GRID



# Why?

---



# Flattening !

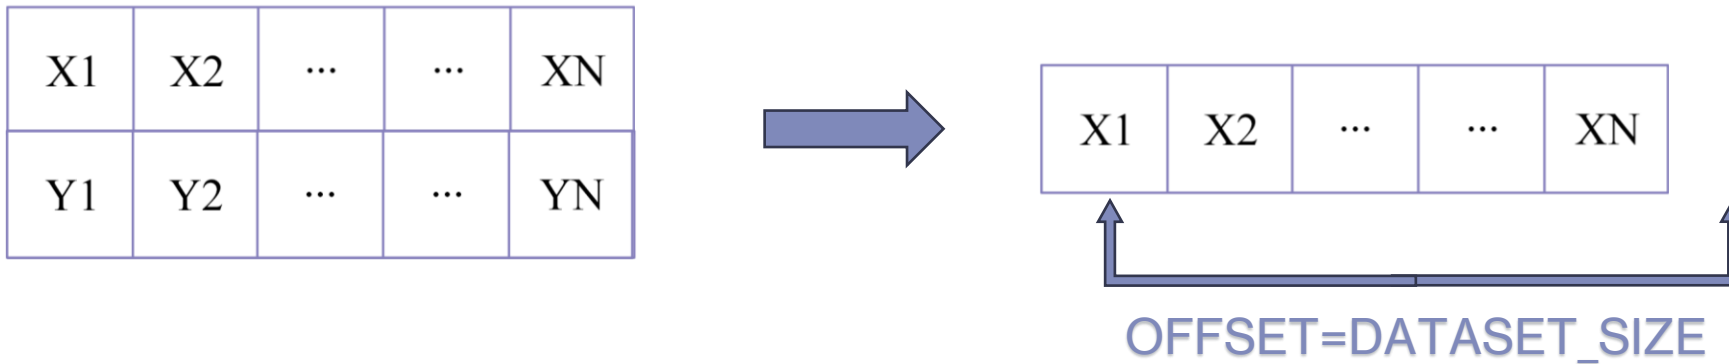
On gpu, if we are operating on bi-dimensional data, if is possible, is highly recommended to perform a flattening and operate on one dimension data structure.

Our first implementation already used a data structure that exploited a sort of flattening.

```
typedef struct {  
    double coordinate[3];  
} Point;  
Point* points = (Point*)malloc(DATASET_SIZE_HOST * sizeof(Point));  
Point* d_points;  
cudaMalloc((void**)&d_points, DATASET_SIZE_HOST * sizeof(Point));
```

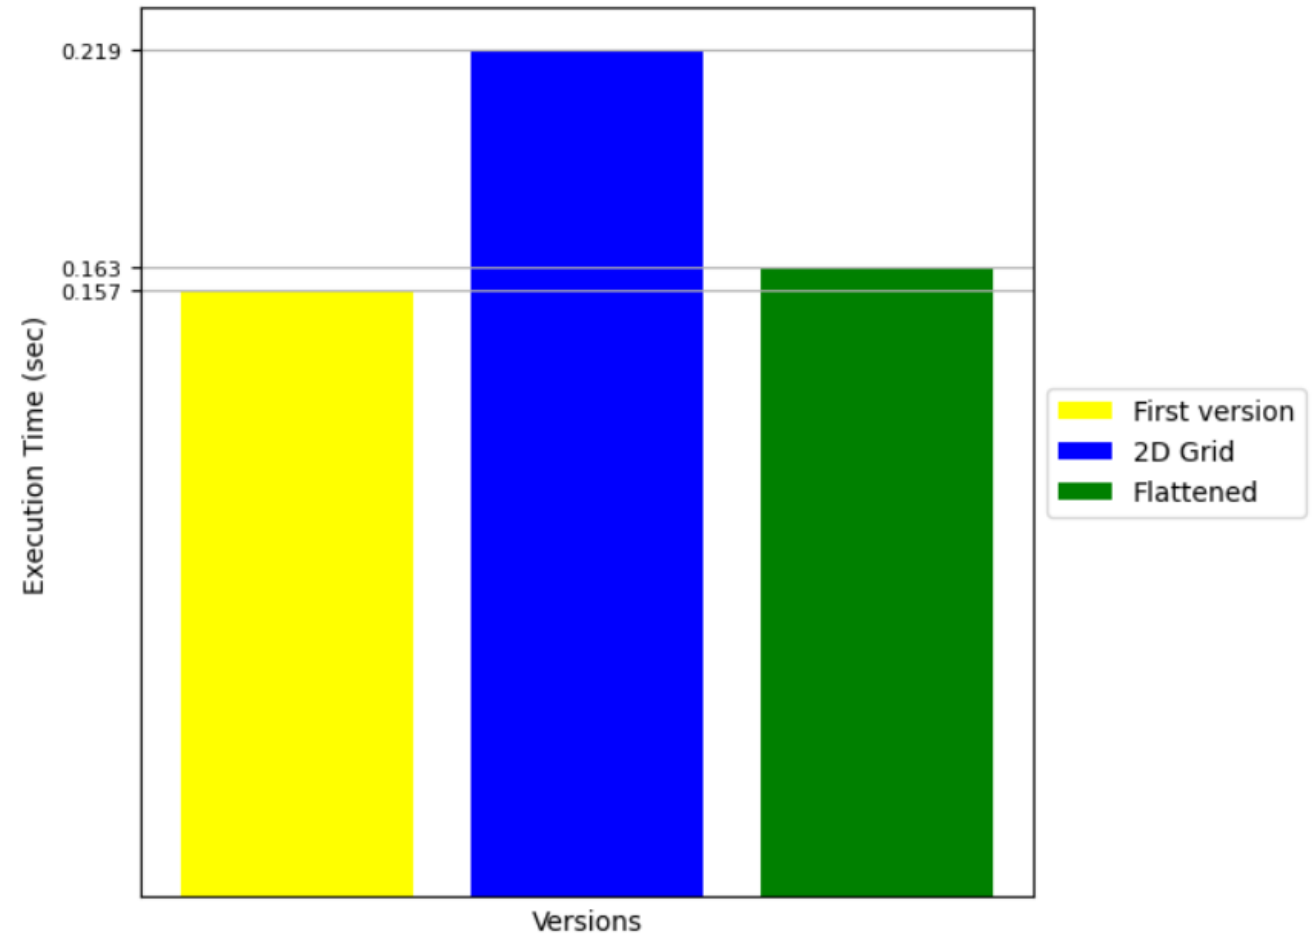
# Flattening !

To prove this supposition we used another data structure, a 2D matrix flattened to an array.



```
double* d_pointsArray;  
cudaMalloc((void**)&d_pointsArray, DIM * DATASET_SIZE * sizeof(double));
```

# Comparison of execution time on 10 millions points between initial version, 2D Grid and Flattened version




# GPU – Optimization

Reducing the number of calls on atomic operations: is there any way?

assignPointsToClusterKernel

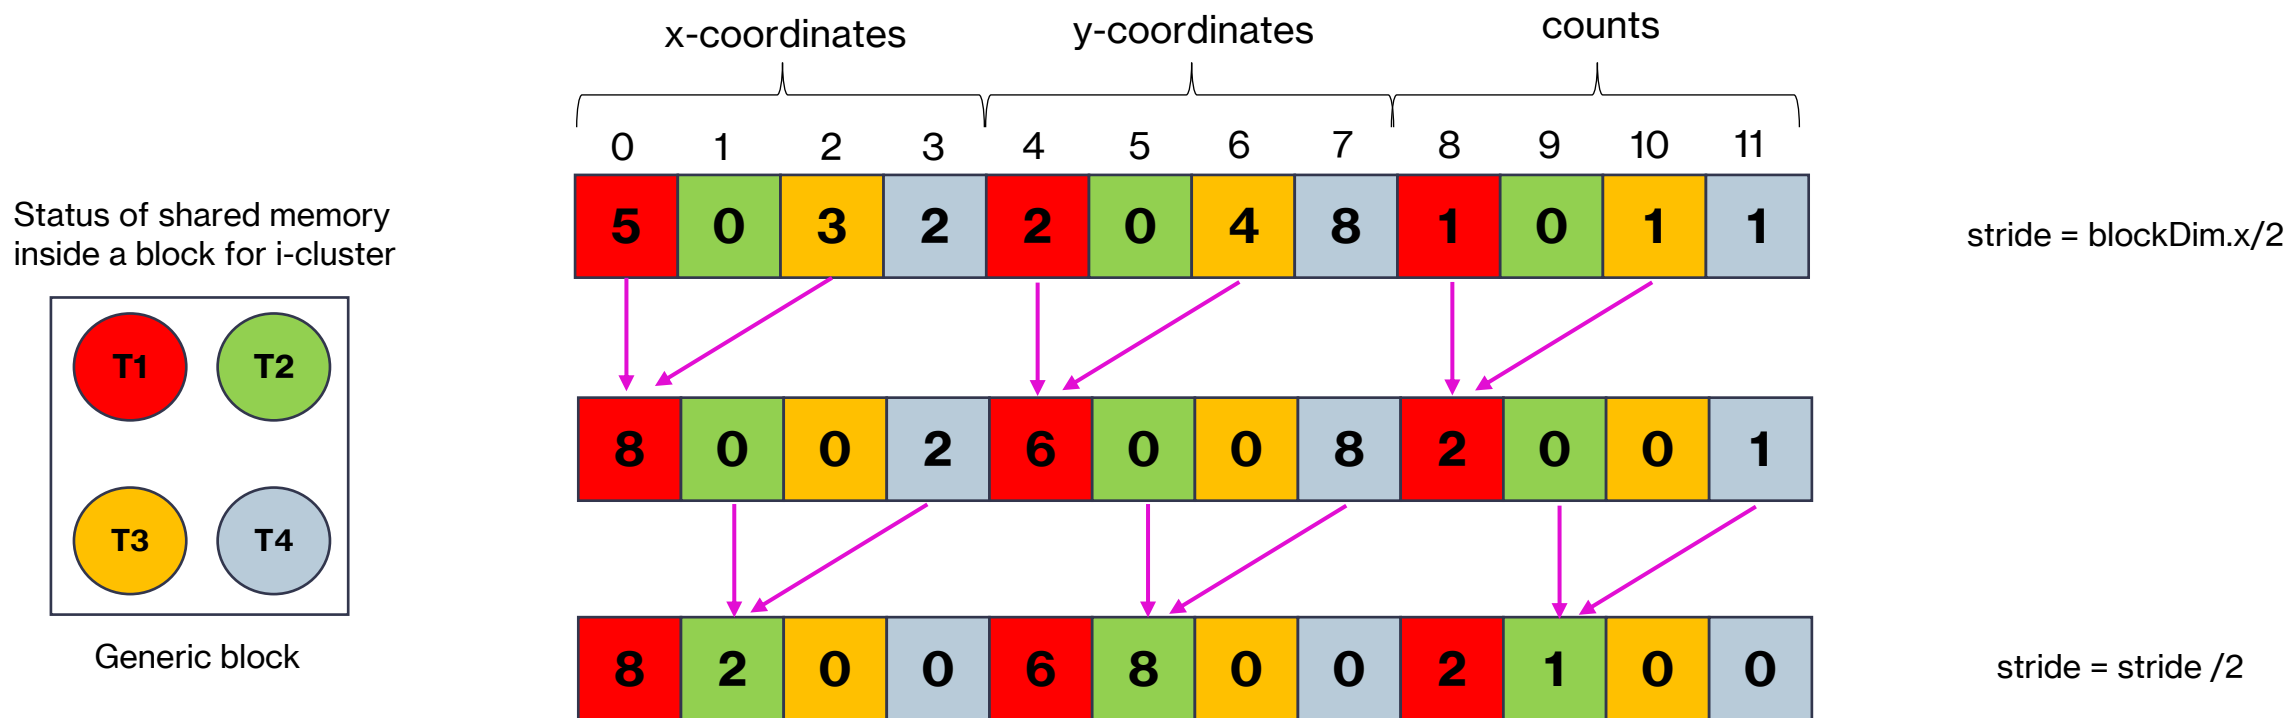
```
// Update sums and counts
for (int dim = 0; dim < DIM; dim++) {
    atomicAdd(&sums[minIndex].coordinate[dim], points[idx].coordinate[dim]);
}
atomicAdd(&counts[minIndex], 1);
```

How many times?  $\text{DATASET\_SIZE} \times (\text{DIM} + 1)$   In our case:  $\text{DIM}=2$  and  $\text{DATASET\_SIZE} = 10\,000\,000$   
30 000 000 calls per iteration

Can we do better exploiting our GPU architecture?

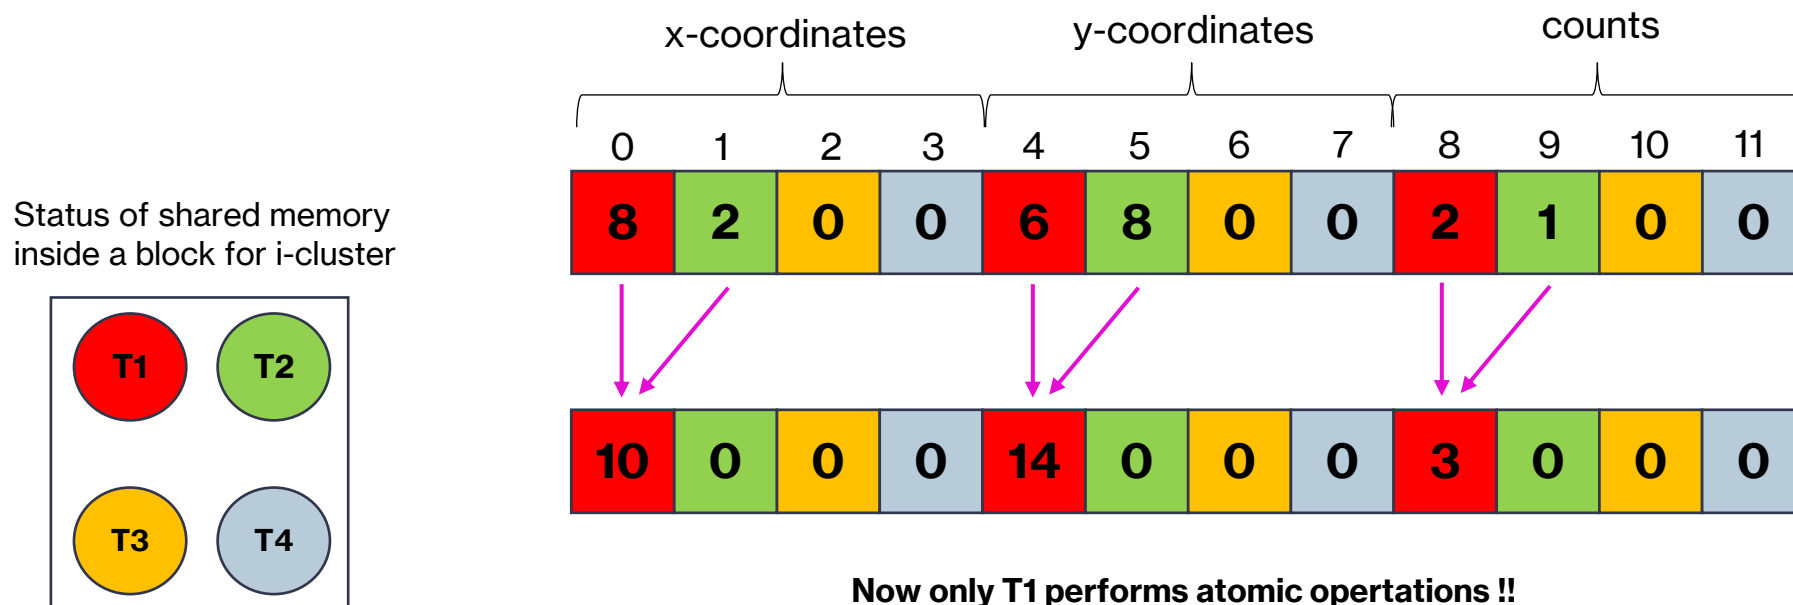
# GPU – Optimization

Rewrite `assignPointsToClusterKernel` in order to perform for each cluster a PARALLEL - TREE - REDUCTION exploiting the shared memory in each block



# GPU – Optimization

Rewrite `assignPointsToClusterKernel` in order to perform for each cluster a PARALLEL - TREE – REDUCTION exploiting the shared memory of each block



**A BLOCK NEVER  
HANDLES ONLY  
4 THREADS!!**

This is just an example: let see the impact on real data



# GPU – Optimization

Reducing the number of calls on atomic operations

- Initial version:  $\text{DATASET\_SIZE} \times (\text{DIM} + 1) \rightarrow 30\,000\,000$  calls per iteration
- Opt. version:  $\text{NR\_BLOCKS} \times (\text{DIM} + 1) \times K \rightarrow \sim 150\,000$  calls per iteration

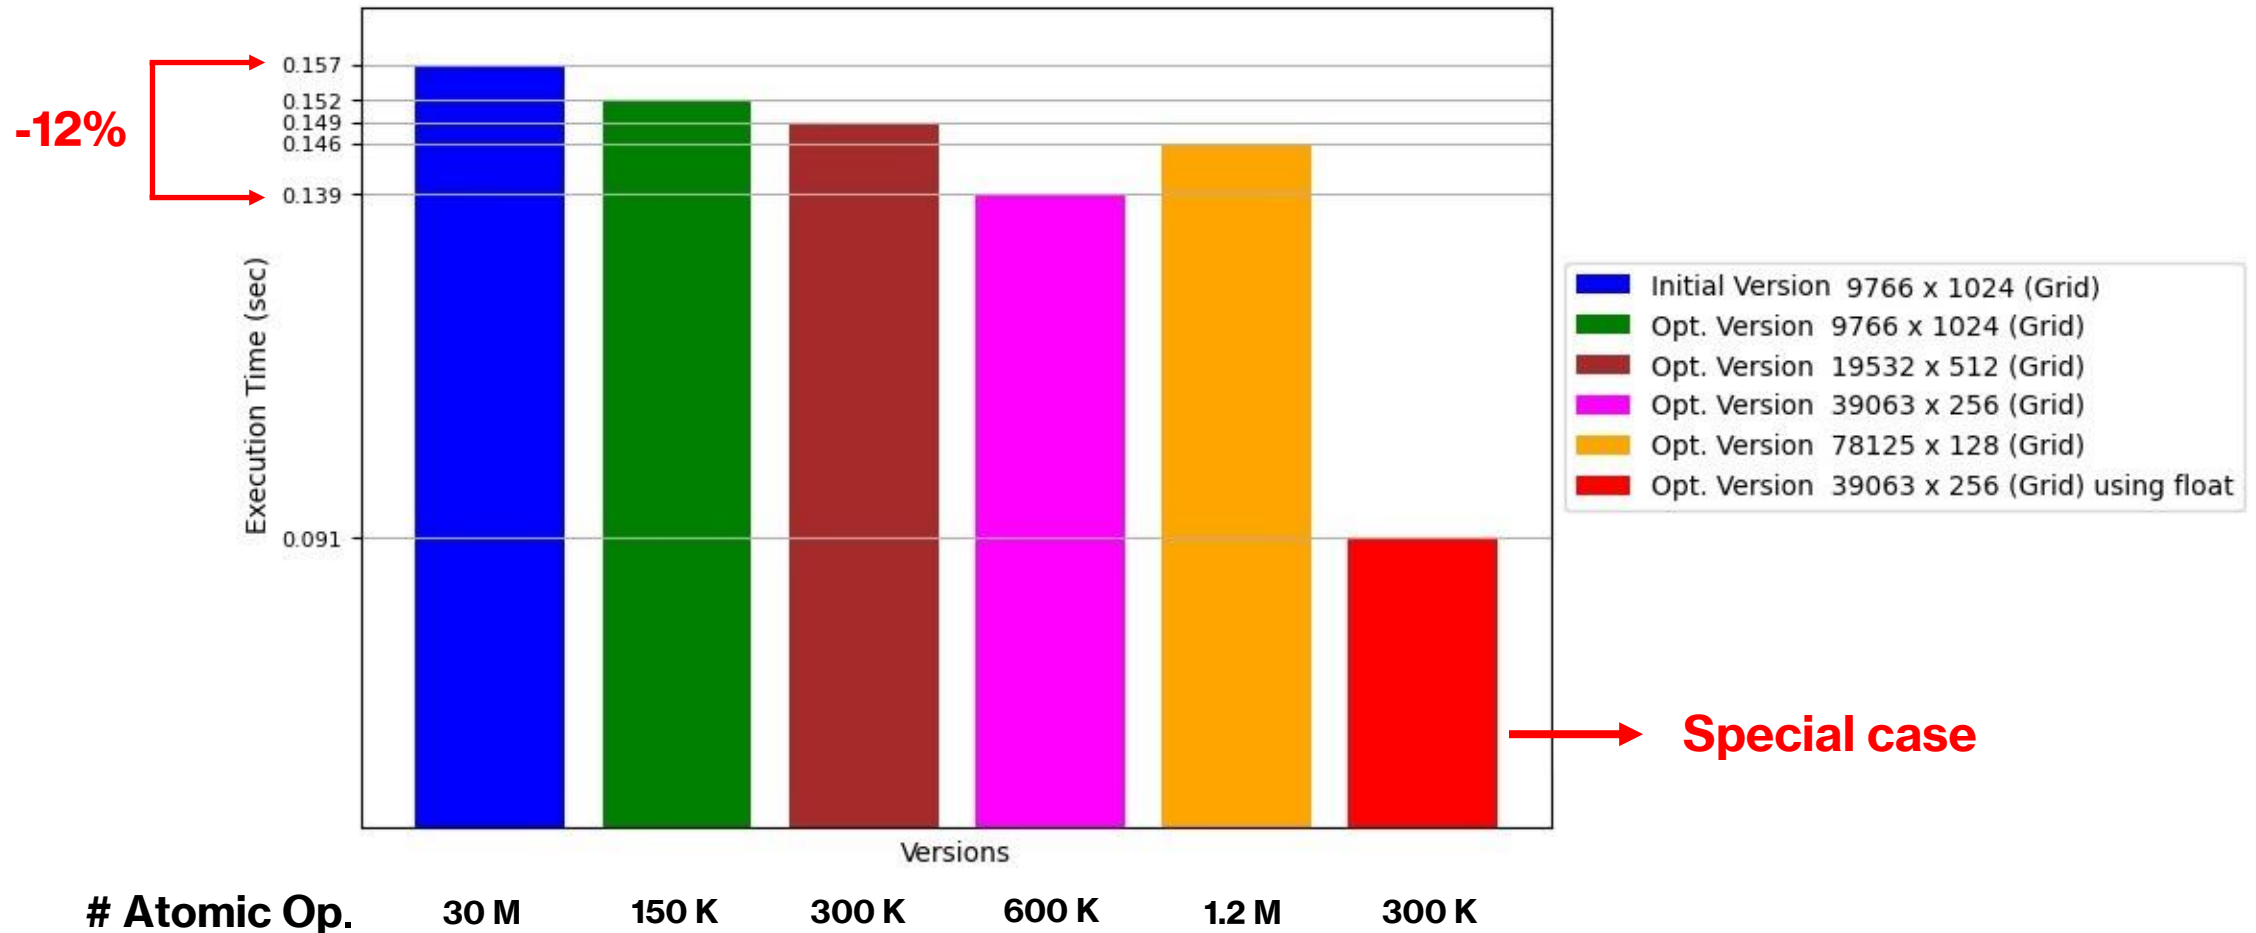
**DIM=2**  
**DATASET\_SIZE = 10 000 000**  
**NR\_BLOCKS=9766**  
**THREADS=1024**  
**K(# of clusters)=5**

PAY ATTENTION: the opt. version introduces additional synchronizations among threads in the same block

```
__syncthreads();
```

We have to find a trade-off between # of blocks and # of threads

# GPU – Optimization



# GPU – Optimization



v.2019

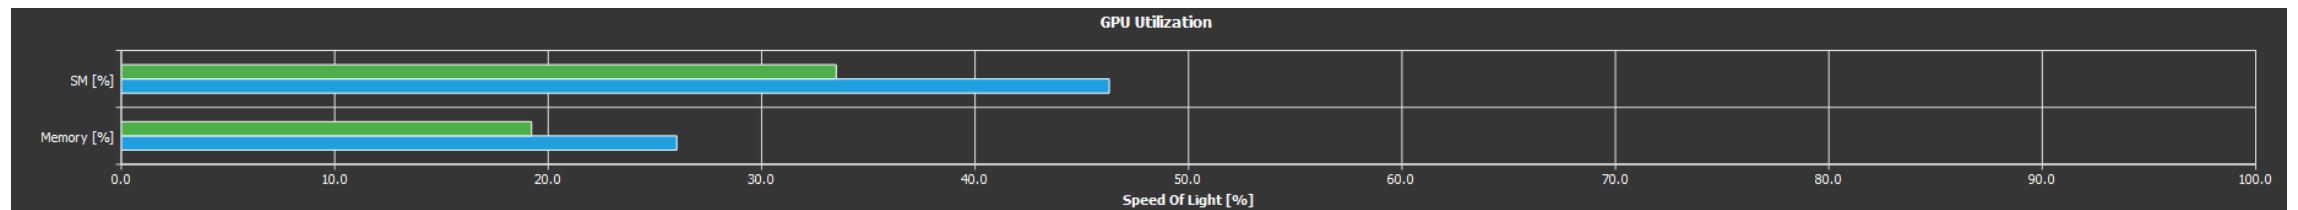


Without parallel tree



With parallel tree

SOL SM [%]	46,29	(+38,18%)
SOL Memory [%]	26,03	(+35,40%)
SOL TEX [%]	18,10	(+16,06%)
SOL L2 [%]	21,69	(+12,83%)
SOL FB [%]	10,11	(+15,78%)



# GPU – Optimization



NVIDIA  
Nsight™  
Compute

v.2019

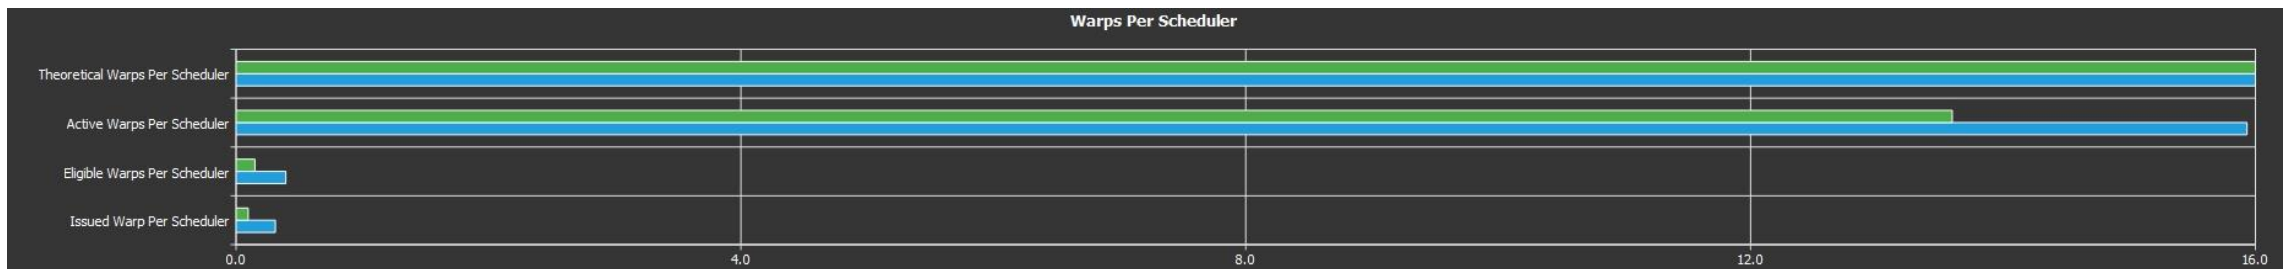


Without parallel tree



With parallel tree

Active Warps Per Scheduler [warp/cycle]	15,93	(+17,18%)
Eligible Warps Per Scheduler [warp/cycle]	0,39	(+163,32%)
Issued Warp Per Scheduler [issue/cycle]	0,31	(+223,54%)
Instructions Per Active Issue Slot [inst/issue]	1,20	(+12,64%)
No Eligible [%]	68,91	(-23,65%)
One or More Eligible [%]	31,14	(+223,54%)



# GPU – Optimization



v.2019

```
__syncthreads();
```

● Without parallel tree

● With parallel tree



# Optimizations summary

For CPU Parallel version:

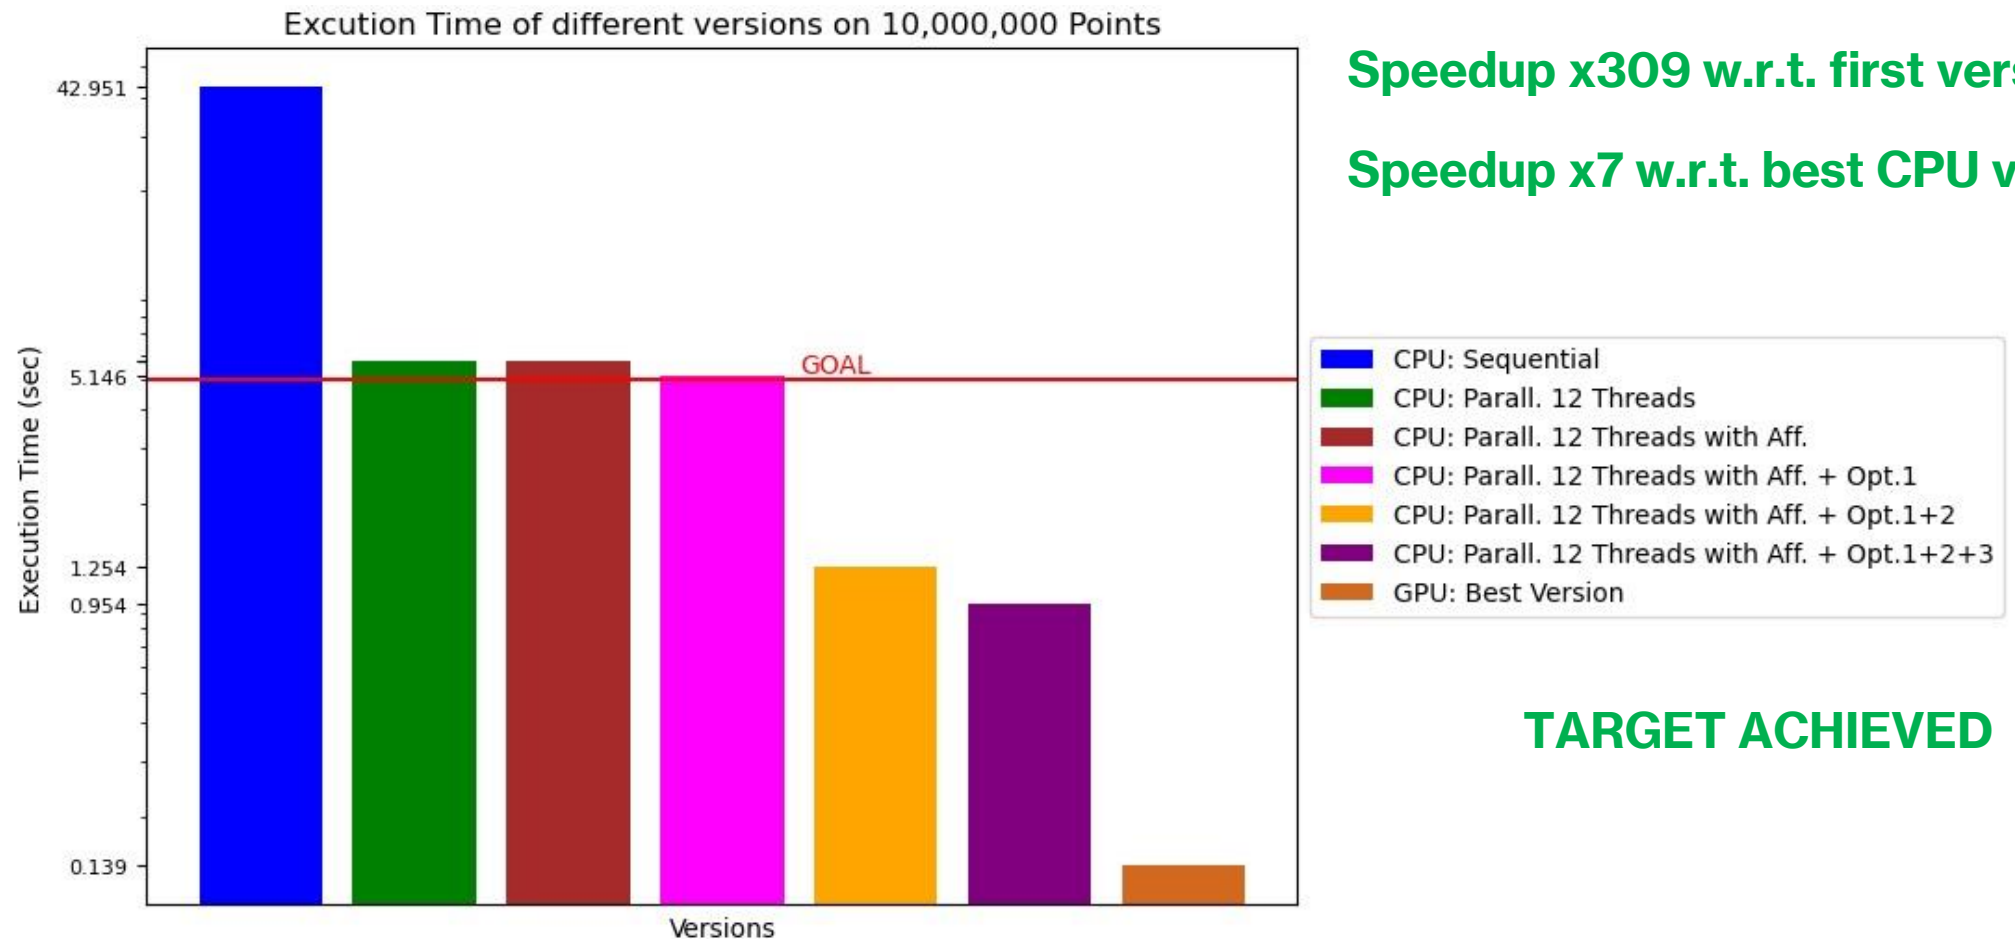
1. From a dynamic list of points, to an array of points [Opt.1];
2. Switching from usage of `pow()` function, to executing directly the square of the number ( $N^2 = N \times N$ ) [Opt.2];
3. Find the minimum distance between to points, without using the square root [Opt.3];

For GPU version:

1. Usage of parallel tree reduction technique;



# Conclusions



Speedup x309 w.r.t. first version

Speedup x7 w.r.t. best CPU version

**TARGET ACHIEVED**

# Thank you

---